



# vLLM Office Hours #39

**Special Topic:** Intro to batch invariant in vLLM

*January 8, 2026*



**Michael Goin**

Principal Engineer, Red Hat  
vLLM Committer



**Saša Zelenović**

Developer Marketing and Advocacy Manager, Red Hat



**Wentao Ye**

Machine Learning Engineer, Red Hat  
vLLM Committer



**Bram Wasti**

Software Engineer, Meta

# What's new in the past two weeks?

## vLLM Project Update

- ▶ vLLM-Omni v0.12.Orc1
- ▶ vLLM Semantic Router v0.1
- ▶ vLLM v0.13.0

## Upcoming vLLM Office Hours Sessions

- ▶ [Dec 18] [vLLM 2025 Retrospective & 2026 Roadmap](#)
- ▶ [TODAY] Intro to batch invariant in vLLM
- ▶ [Jan 15] Intro to Speculators, a unified library for building and storing speculative decoding algorithms for LLMs with vLLM
- ▶ [Jan 22] LLM Compressor update
- ▶ [Jan 29] Deep Dive into the vLLM CPU offloading connector

Register for all sessions [here](#).

View previous recordings [here](#).



# Reminder: New vLLM Website & Events Calendar

vLLM

Event Calendar

View the next two weeks at a glance · Hover for details

Synced · just now

Dec 14 – Dec 27, 2025 Today 1 event in 1 day

SUN	MON	TUE	WED	THU	FRI	SAT
14	15	16	17	18	19	20
—	—	—	—	1 event	—	—
SUN	MON	TUE	WED	THU	FRI	SAT
21	22	23	24	25	26	27
—	—	—	—	—	—	—

Meetup • Office Hours • Release • News • Webinar • Conference

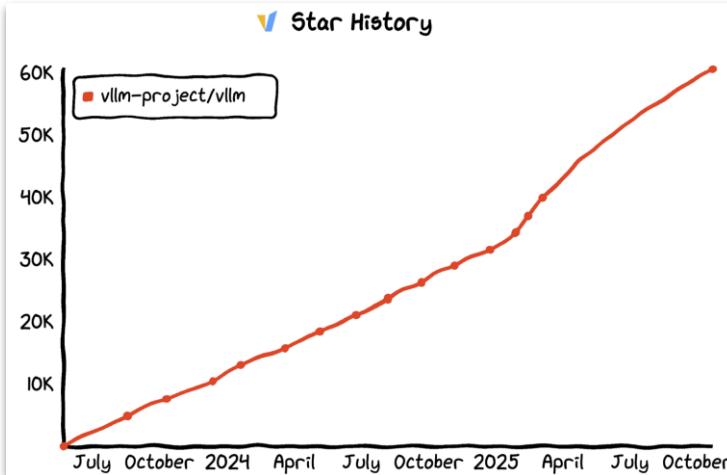
Upcoming Events 3

- Dec 18, 2025 · 3:00 PM America/New\_York  
vLLM Office Hours #38
- Jan 8, 2026 · 2:00 PM America/New\_York  
vLLM Office Hours #39
- Jan 15, 2026 · 2:00 PM America/New\_York  
vLLM Office Hours #40

▶ Head over to [vllm.ai/events](https://vllm.ai/events) to see upcoming office hours, meetups, conferences, etc.

# What is vLLM?

The High-Throughput and Memory-Efficient inference and serving engine for LLMs



## Most Popular LLM Serving Engine

- 65K+ GitHub stars, 800+ PRs/month
- 500K++ GPUs deployed 24/7
- 2K+ contributors, 10K+ members in [slack.vllm.ai](https://slack.vllm.ai)

>2000 Contributors from >50 major companies



<https://github.com/vllm-project/vllm>

```
$ uv pip install vllm --torch-backend=auto  
$ vllm serve deepseek-ai/DeepSeek-V3.1 -tp 8
```

## Broad Model Support (>100 arches)



## Flexible Device Parallelism

Tensor, Pipeline, Expert, Data, Context Parallel  
Disagg Prefill/Decode, Disagg Encoder

## Wide Hardware Support



## Diverse Project Ecosystem



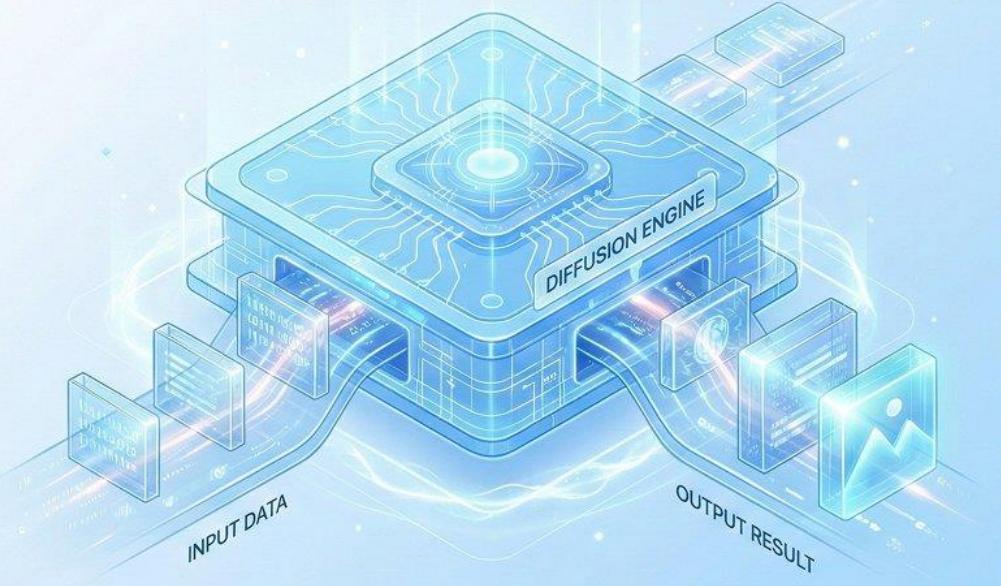
# vLLM-Omni v0.12.Orc1 Pre-Release

<https://github.com/vllm-project/vllm-omni/releases/tag/v0.12.Orc1>

<https://docs.vllm.ai/projects/vllm-omni/en/latest/>

## vLLM-Omni v0.12.Orc1

Production-Grade Multimodal Serving



The diagram illustrates the vLLM-Omni architecture. It features a central "DIFFUSION ENGINE" represented by a blue cube with a circular interface on top. This engine is connected to a "INPUT DATA" source at the bottom left and an "OUTPUT RESULT" destination at the bottom right. The entire process is depicted within a glowing blue rectangular frame.

### Release Notes

- Video** Wan2.2
- Image** Qwen
- Speed** TeaCache

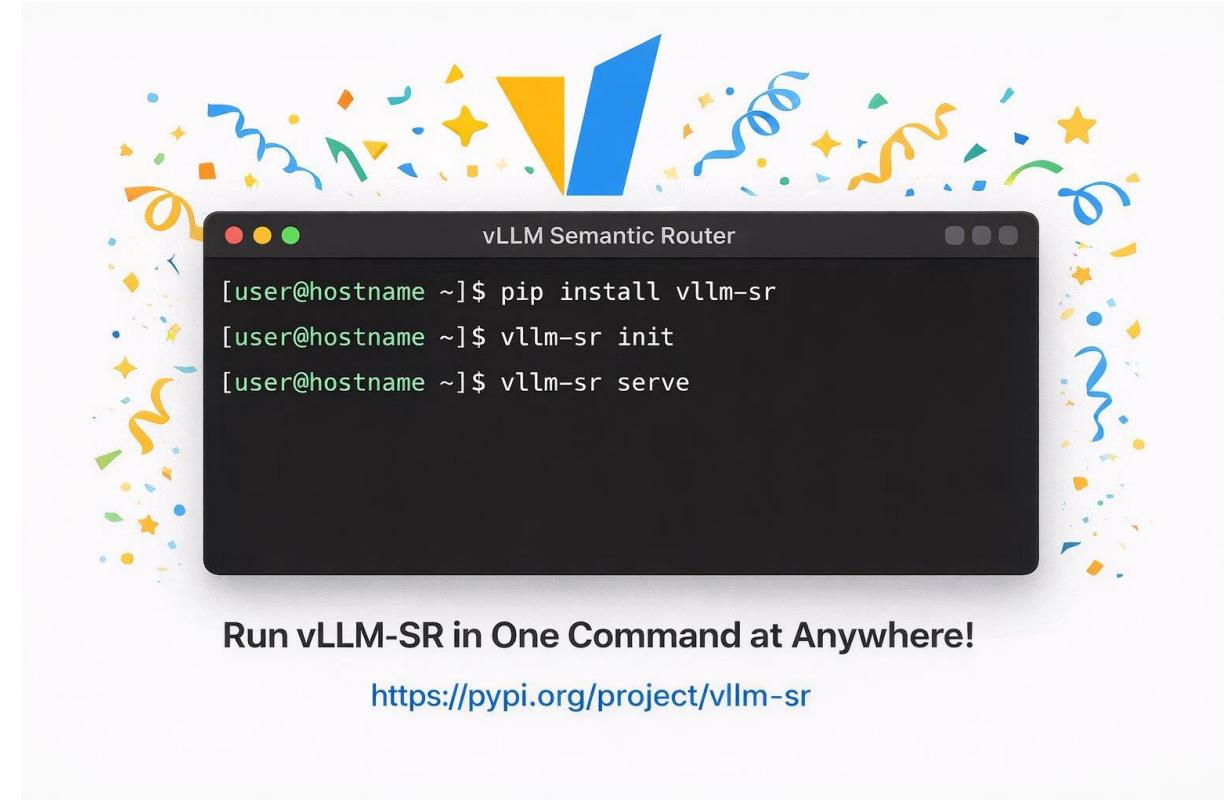
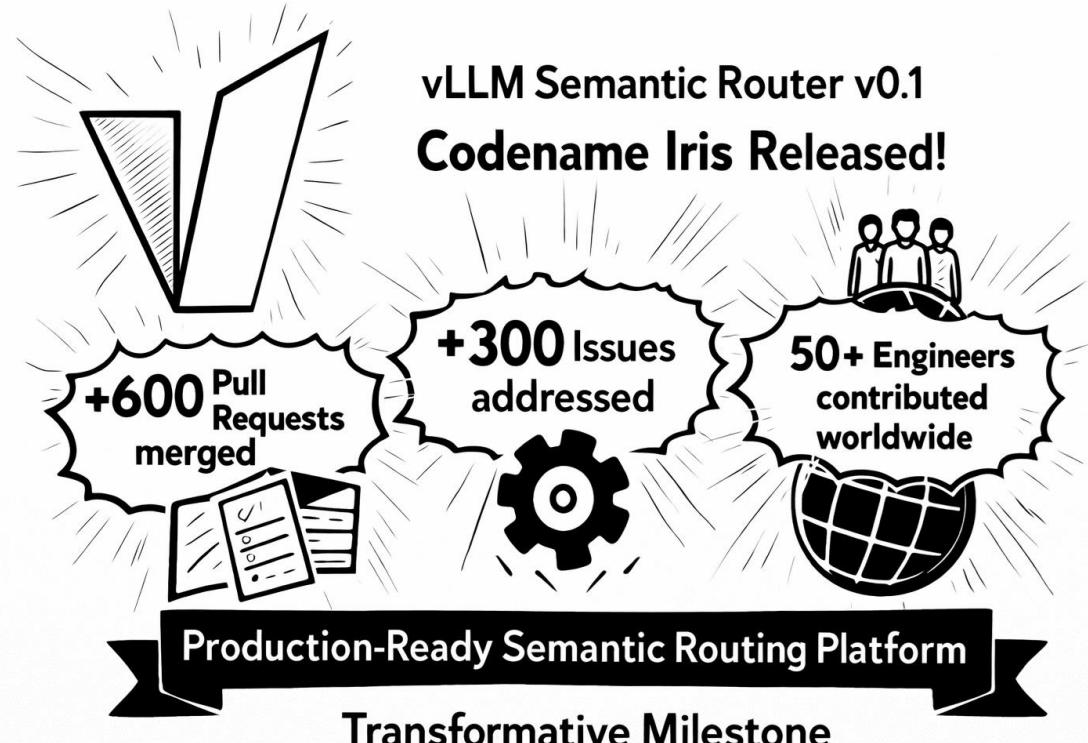
- ✓ Diffusion Engine Overhaul
- ✓ OpenAI Compatible API
- ✓ TeaCache & Cache-DiT
- ✓ AMD ROCm Support

187 Commits    45 Contributors



# vLLM Semantic Router v0.1 "Iris" Release

<https://blog.vllm.ai/2026/01/05/vllm-sr-iris.html>



Thanks to the many clean, from-scratch implementations of vLLM!

<https://github.com/GeeeekExplorer/nano-vllm>

<https://github.com/Wenyueh/MinivLLM>

<https://github.com/skyzh/tiny-llm>



# What's new in vLLM v0.13.0

🎉 442 commits from 207 contributors, including 61 new contributors! 🎉

## Model Support

- ▶ New models: BAGEL (AR only) ([#28439](#)), AudioFlamingo3 ([#30539](#)), JAIS 2 ([#30188](#)), Nemotron latent MoE ([#30203](#)).
- ▶ Tool parsers: DeepSeek-V3.2 ([#29848](#)), Gigachat 3 ([#29905](#)), Holo2 reasoning ([#30048](#)).
- ▶ Model enhancements: Qwen3-VL embeddings ([#30037](#)), Qwen3-VL EVS (Efficient Video Sampling) ([#29752](#)), DeepSeek V3.2 proper drop\_thinking logic ([#30490](#)) and top-k fix ([#27568](#)).
- ▶ Task expansion: Automatic TokenClassification model conversion ([#30666](#)), Ultravox v0.7 transformer projector ([#30089](#)).
- ▶ Quantization: BitsAndBytes for Qwen3-Omni-MoE ([#29896](#)).
- ▶ Speculative decoding: Eagle/Eagle3 Transformers backend ([#30340](#)), Mamba selective\_state\_update spec decode ([#29488](#)).

## Engine Core

- ▶ Compilation via `compile_ranges` for selective kernel compilation ([#24252](#)).
- ▶ Prefix caching: xxHash high-performance hash option ([#29163](#)).
- ▶ Attention: PrefixLM support for FlexAttention ([#27938](#)) and TritonAttention ([#30386](#)), CUDA graphs for 3D Triton attention ([#28306](#)).
- ▶ Batch invariance: FA2 and LoRA batch-invariant support ([#30018](#)), TRITON\_MLA without prefix-caching ([#29125](#)).
- ▶ Pooling: Chunked prefill for ALL pooling tasks ([#27145](#)), multi-vector retrieval API ([#26686](#)).
- ▶ Model Runner V2: Min-p ([#30171](#)), NaN detection in logits ([#30187](#)).
- ▶ Speculative decoding: Medusa GPU-CPU sync avoidance ([#29723](#)), async spec-decode improvements ([#29624](#)).
- ▶ Whisper: ~3x speedup vs v0.12.0, Encoder batching ([#29421](#)), FULL\_DECODE\_ONLY CUDA graph ([#30072](#)), CPU support ([#30062](#)).
- ▶ Performance: Fused blockwise quant RMS norm ([#27883](#)), MoE LoRA loading reduction ([#30243](#)), encoder cache optimization ([#30475](#)), CPU KV offloading streams ([#29013](#)).

# What's new in vLLM v0.13.0



442 commits from 207 contributors, including 61 new contributors!



## Hardware & Performance

- ▶ NVIDIA Blackwell Ultra SM103 (GB300) support ([#30484](#)).
- ▶ Several DeepSeek/Kimi optimizations:
  - DeepEP High-Throughput CUDA graph enabled by default: 5.3% throughput, 4.4% TTFT improvement ([#29558](#))
  - DeepGEMM fused layout: 10.7% TTFT improvement ([#29546](#))
  - DeepGEMM experts init: 3.9% TTFT improvement ([#30494](#))
  - group\_topk kernel: 2% throughput improvement ([#30159](#))
  - Sparse prefill kernel for DeepSeek-V3.2 FP8 KV-cache ([#27532](#))
  - MLA FP8 Quant ([#29795](#)), broadcast k\_nope/k\_pe ([#29710](#))
- ▶ CPU: Whisper support ([#30062](#)), Arm vectorized exp ([#30068](#)), x86 CPU wheel pipeline ([#28848](#)).
- ▶ AMD ROCm: Aiter quantization kernels ([#25552](#)), torch.compile layernorm/silu + FP8 quant ([#25693](#)), Triton ScaledMM fallback ([#26668](#)), MXFP4 w4a4 inference ([#29775](#)).
- ▶ Intel XPU: wNa16 compressed tensors ([#29484](#)).
- ▶ Build: CUDA 13 aarch64 wheels ([#30341](#)), Docker kernel build stage ([#29452](#)), Ascend NPU Docker ([#30015](#)).

## Large Scale Serving & Disaggregated Prefill/Decode

- ▶ KV connectors: Mooncake Transfer Engine ([#24718](#)), cache reset via /reset\_prefix\_cache ([#27170](#)), KV events ([#28309](#)), failure recovery config ([#26813](#)).
- ▶ NIXL: Compatibility checking in handshake ([#29503](#)), large batch proxy support ([#28782](#)).
- ▶ EPLB: NVFP4 support ([#29804](#)), algorithm abstraction ([#26471](#)).
- ▶ Multi-node: External launcher mode ([#29833](#)).
- ▶ Hybrid allocator: Optional KV connector integration ([#29805](#)).
- ▶ Performance: silu\_mul\_per\_token\_group\_quant\_fp8 kernel for DP/EP ([#29470](#)).

# What's new in vLLM v0.13.0



442 commits from 207 contributors, including 61 new contributors!



## API & Frontend

- ▶ Responses API: MCP type infrastructure ([#30054](#)), Browser/Container MCP tools ([#29989](#)), full MCP Python loop ([#29798](#)), extra body parameters ([#30532](#)).
- ▶ Configuration: AttentionConfig replaces `VLLM_ATTENTION_BACKEND` env var ([#26315](#)).
- ▶ Chat templates: DeepSeek-V3.2 ([#29837](#)), DeepSeek-V3.2 developer tools ([#30040](#)).
- ▶ Anthropic API: Streaming fixes ([#29971](#), [#30266](#)).
- ▶ Embeddings: Binary format with `encoding_format=bytes_only` ([#30249](#)), multiple image/audio per request ([#29988](#)), `tokenization_kwarg` override ([#29794](#)).
- ▶ Metrics: Prefill KV compute metric excluding cached tokens ([#30189](#)).
- ▶ Profiling: Layer-wise NVTX ([#29990](#)), profiling CLI config ([#29912](#)).
- ▶ UX: Better OOM errors ([#28051](#)), ModelConfig validation ([#30213](#)), distributed executor errors ([#30140](#)).

## Breaking Changes & Deprecations

This release includes deprecation removals, PassConfig flag renames, and attention configuration changes from environment variables to CLI arguments. **Please review the breaking changes section carefully before upgrading.**

- ▶ PassConfig flags renamed per RFC [#27995](#) ([#29646](#))
- ▶ Attention env vars → CLI args: `VLLM_ATTENTION_BACKEND` replaced with `--attention-backend` ([#26315](#))
- ▶ Removed `-o.xx` flag ([#29991](#))
- ▶ Removed deprecated plugin/compilation fields ([#30396](#))
- ▶ Removed deprecated task, seed, MM settings ([#30397](#))
- ▶ Removed `embed_input_ids/embed_multimodal` fallbacks ([#30458](#))
- ▶ Removed tokenizer setter ([#30400](#))
- ▶ Deprecations: `merge_by_field_config` ([#30035](#), [#30170](#)), `--convert_reward` → `--convert_embed` ([#30463](#))

Thank you to the over 2000 contributors!

<https://blog.vllm.ai/2025/12/15/vllm-epd.html>

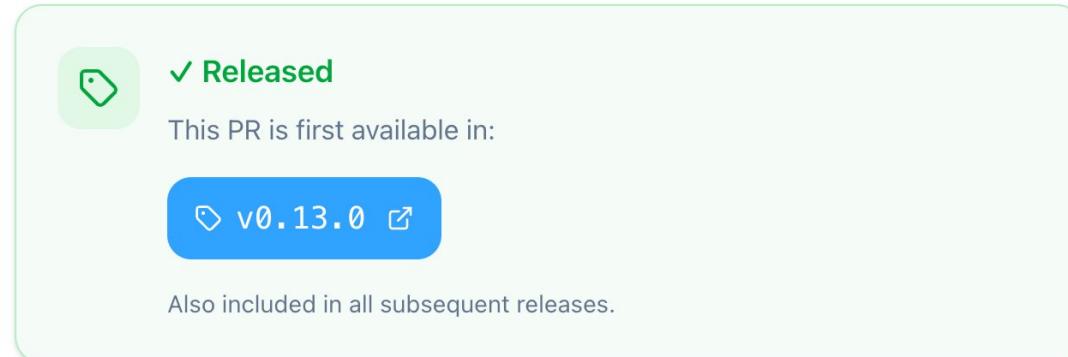


# A thank you gift: PR Release Finder

Ever wondered "Which release first included my PR?"

Just enter your PR number or URL to track your code's journey into production.

[vllm.ai/pr-lookup](https://vllm.ai/pr-lookup)



Today's special topic:

# Intro to batch invariance in vLLM



**Wentao Ye**  
Machine Learning Engineer, Red Hat  
vLLM Committer



**Bram Wasti**  
Software Engineer, Meta  
vLLM contributor



# Why does this happen?

Question:

"Let \$A\$, \$B\$, \$C\$, and \$D\$ be points on the hyperbola: .....  
Find the greatest real number that is less than \$BD^2\$ for all such rhombi."



Greedy, Seed=42, BS=32, #GPU=4

Okay, so I have this problem ... perpendicular, but in a square,  
... for all such rhombi is  $\boxed{480}$ .



BF16

Okay, so I have this problem ... perpendicular. Wait, no, hold on,  
... for all rhombi is 960

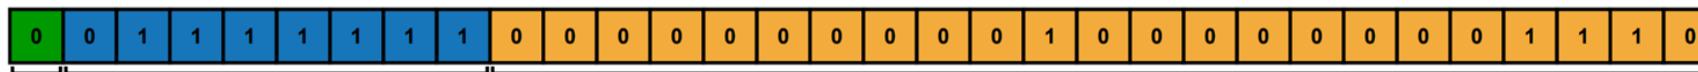


Greedy, Seed=42, BS=8, #GPU=4

Even with a fixed seed and temperature=0?

# Reason 1: Rounding errors of Floating Points

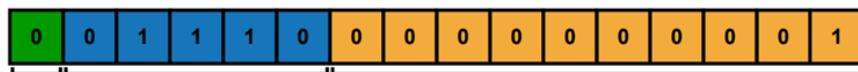
IEEE 754 Single Precision 32-bit Float (IEEE FP32)



Sign: 1 Bit    Exponent: 8 Bits

Mantissa: 23 Bits

IEEE 754 Half Precision 16-bit Float (IEEE FP16)



Sign: 1 Bit    Exponent: 5 Bits

Mantissa: 10 Bits

Google Brain Float (BFloat16 or BF16)



Sign: 1 Bit    Exponent: 8 Bits

Mantissa: 7 Bits

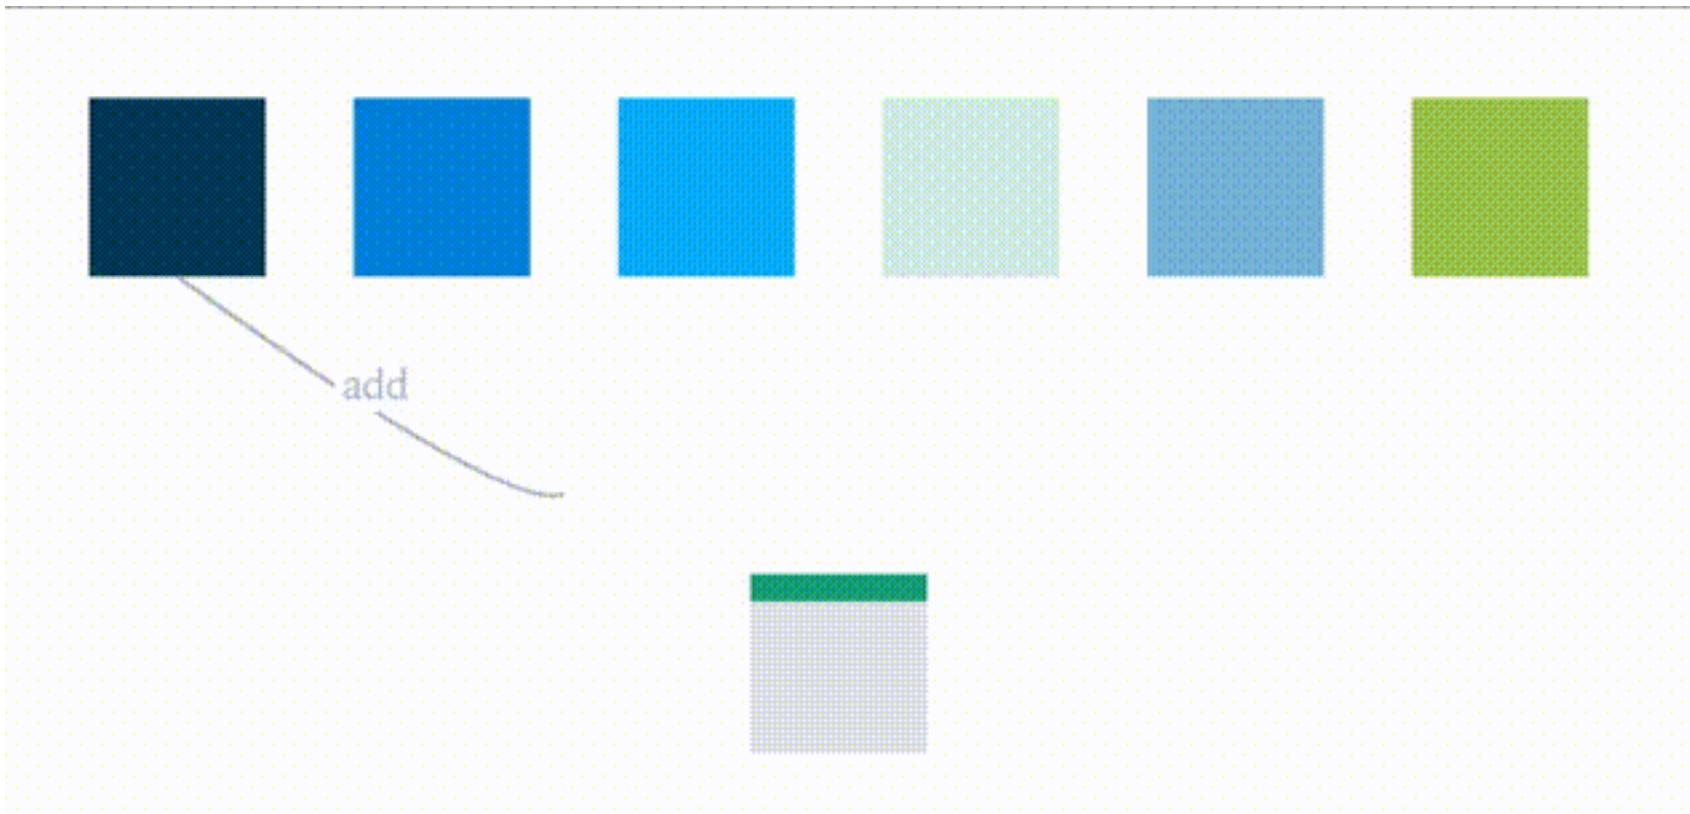
$$\text{Value} = (-1)^{\text{Sign}} \times (1 + \text{Mantissa}) \times 2^{\text{Exponent} - \text{bias}}$$

Precision	Decimal	Rounding Error
FP32	1.00012004375457761	$\approx +4.38e-8$
FP16	1.0	$\approx -0.00012$
BF16	1.0	$\approx -0.00012$

## Reason 2: Non-Associativity

Example	Sum Order	FP32	BF16
$a, b, c = 0.1, -0.1, 0.2$	$a + b + c$	0011111001001100110011001100110011 <b>01</b>	00111110010011 <b>01</b>
	$a + c + b$	0011111001001100110011001100110011 <b>10</b>	00111110010011 <b>10</b>
$a, b, c = 0.0016, 0.0027, 1.0$	$a + b + c$	0011111100000001000110011100111	00111111000000 <b>1</b>
	$a + c + b$	0011111100000001000110011100111	00111111000000 <b>0</b>

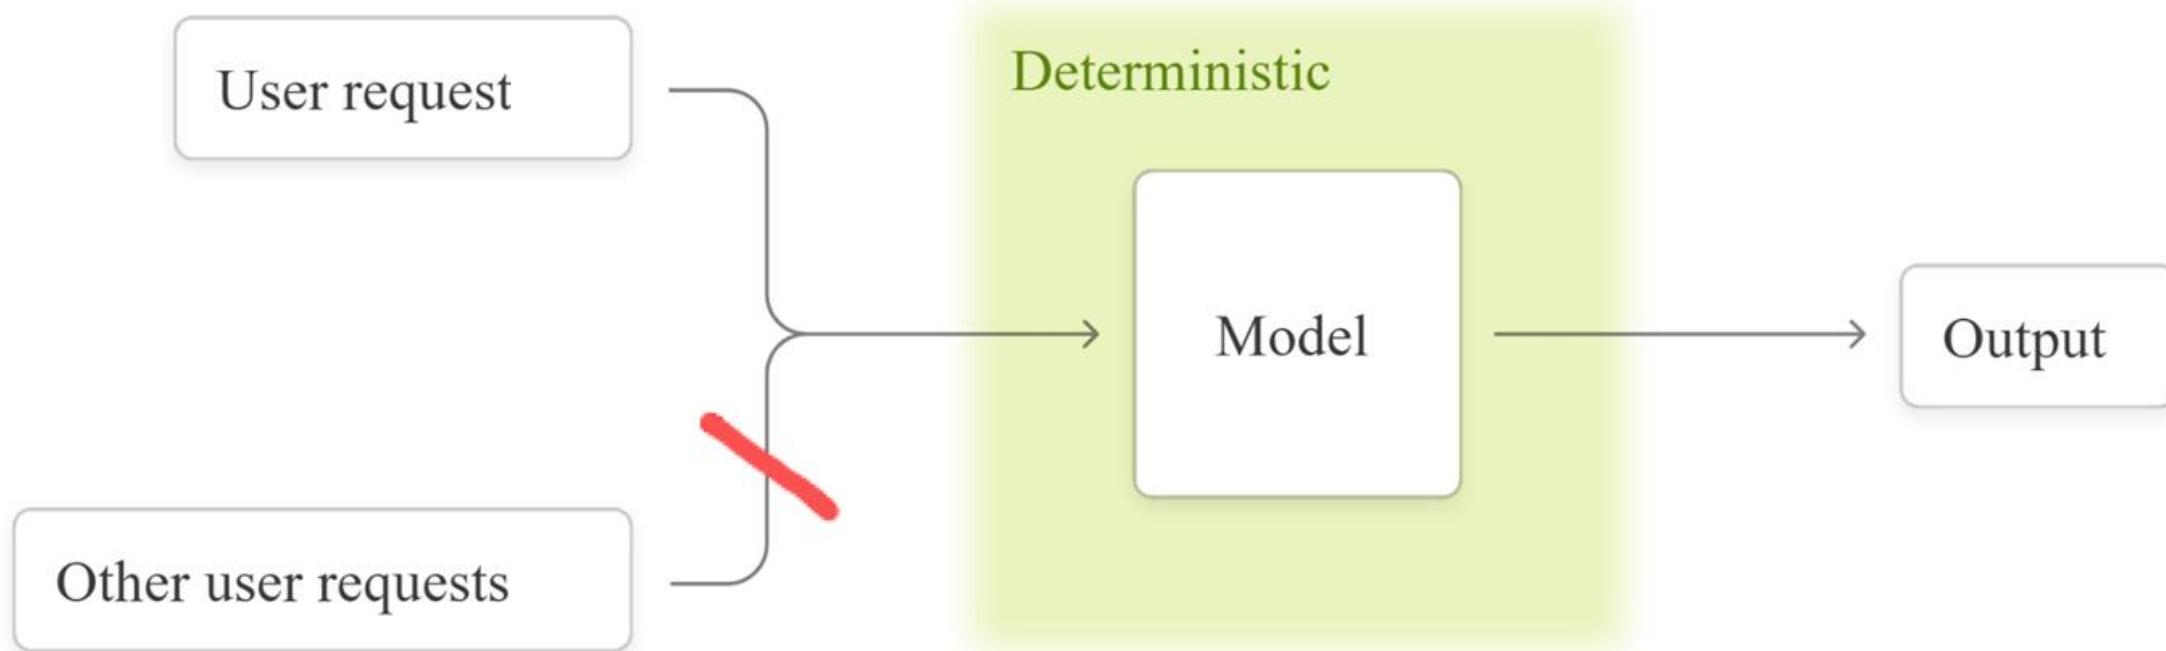
## Reason 1+2 -> Reduction order matters



Sum with atomic add on GPU

Guaranteed for all elements, but the order might change

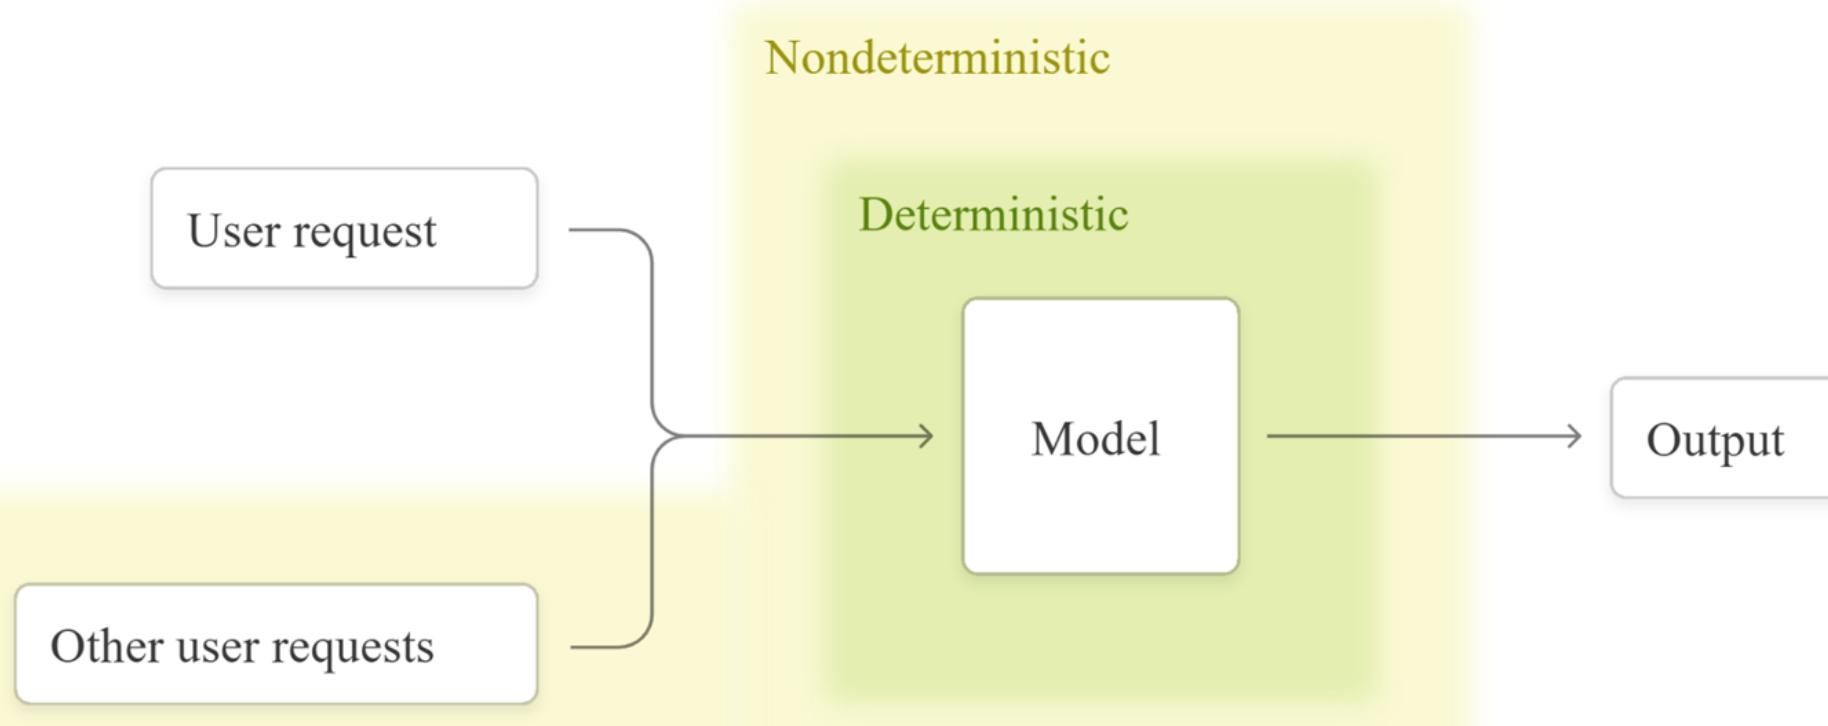
## Reason 3: Continuous Batching



For LLM inference:

No atomic add operator during forward, good for single request

## Reason 3: Continuous Batching



Batching multiple requests together matters

## Reason 3: Continuous Batching

$T_1$	$T_2$	$T_3$	$T_4$	$T_5$	$T_6$	$T_7$	$T_8$
$S_1$	$S_1$	$S_1$	$S_1$	$S_1$	END		
$S_2$	END						
$S_3$	$S_3$	$S_3$	$S_3$	END			
$S_4$	END						

Naive batching, where would  $S_5$  be?

## Reason 3: Continuous Batching

$T_1$	$T_2$	$T_3$	$T_4$	$T_5$	$T_6$	$T_7$	$T_8$
$S_1$	$S_1$	$S_1$	$S_1$				
$S_2$	$S_2$	$S_2$					
$S_3$	$S_3$	$S_3$	$S_3$				
$S_4$	$S_4$	$S_4$	$S_4$	$S_4$			

$T_1$	$T_2$	$T_3$	$T_4$	$T_5$	$T_6$	$T_7$	$T_8$
$S_1$	$S_1$	$S_1$	$S_1$	$S_1$	END	$S_6$	$S_6$
$S_2$	END						
$S_3$	$S_3$	$S_3$	$S_3$	END	$S_5$	$S_5$	$S_5$
$S_4$	$S_4$	$S_4$	$S_4$	$S_4$	$S_4$	END	$S_7$

$S_5$  now starts in a different place

## How to solve this? Batch invariance

```
$ VLLM_BATCH_INVARIANT=0 python invariance_test.py  
Logprob difference: 0.0000052527  
$ VLLM_BATCH_INVARIANT=1 python invariance_test.py  
Logprob difference: 0.0000000000
```

It ensures the output of a model is deterministic and independent of the batch size or the order of requests in a batch (fixed reduction order).

# How to Achieve Batch Invariance?



```
$ VLLM_BATCH_INVARIANT=0 python invariance_test.py
Logprob difference: 0.0000052527
$ VLLM_BATCH_INVARIANT=1 python invariance_test.py
Logprob difference: 0.0000000000
```

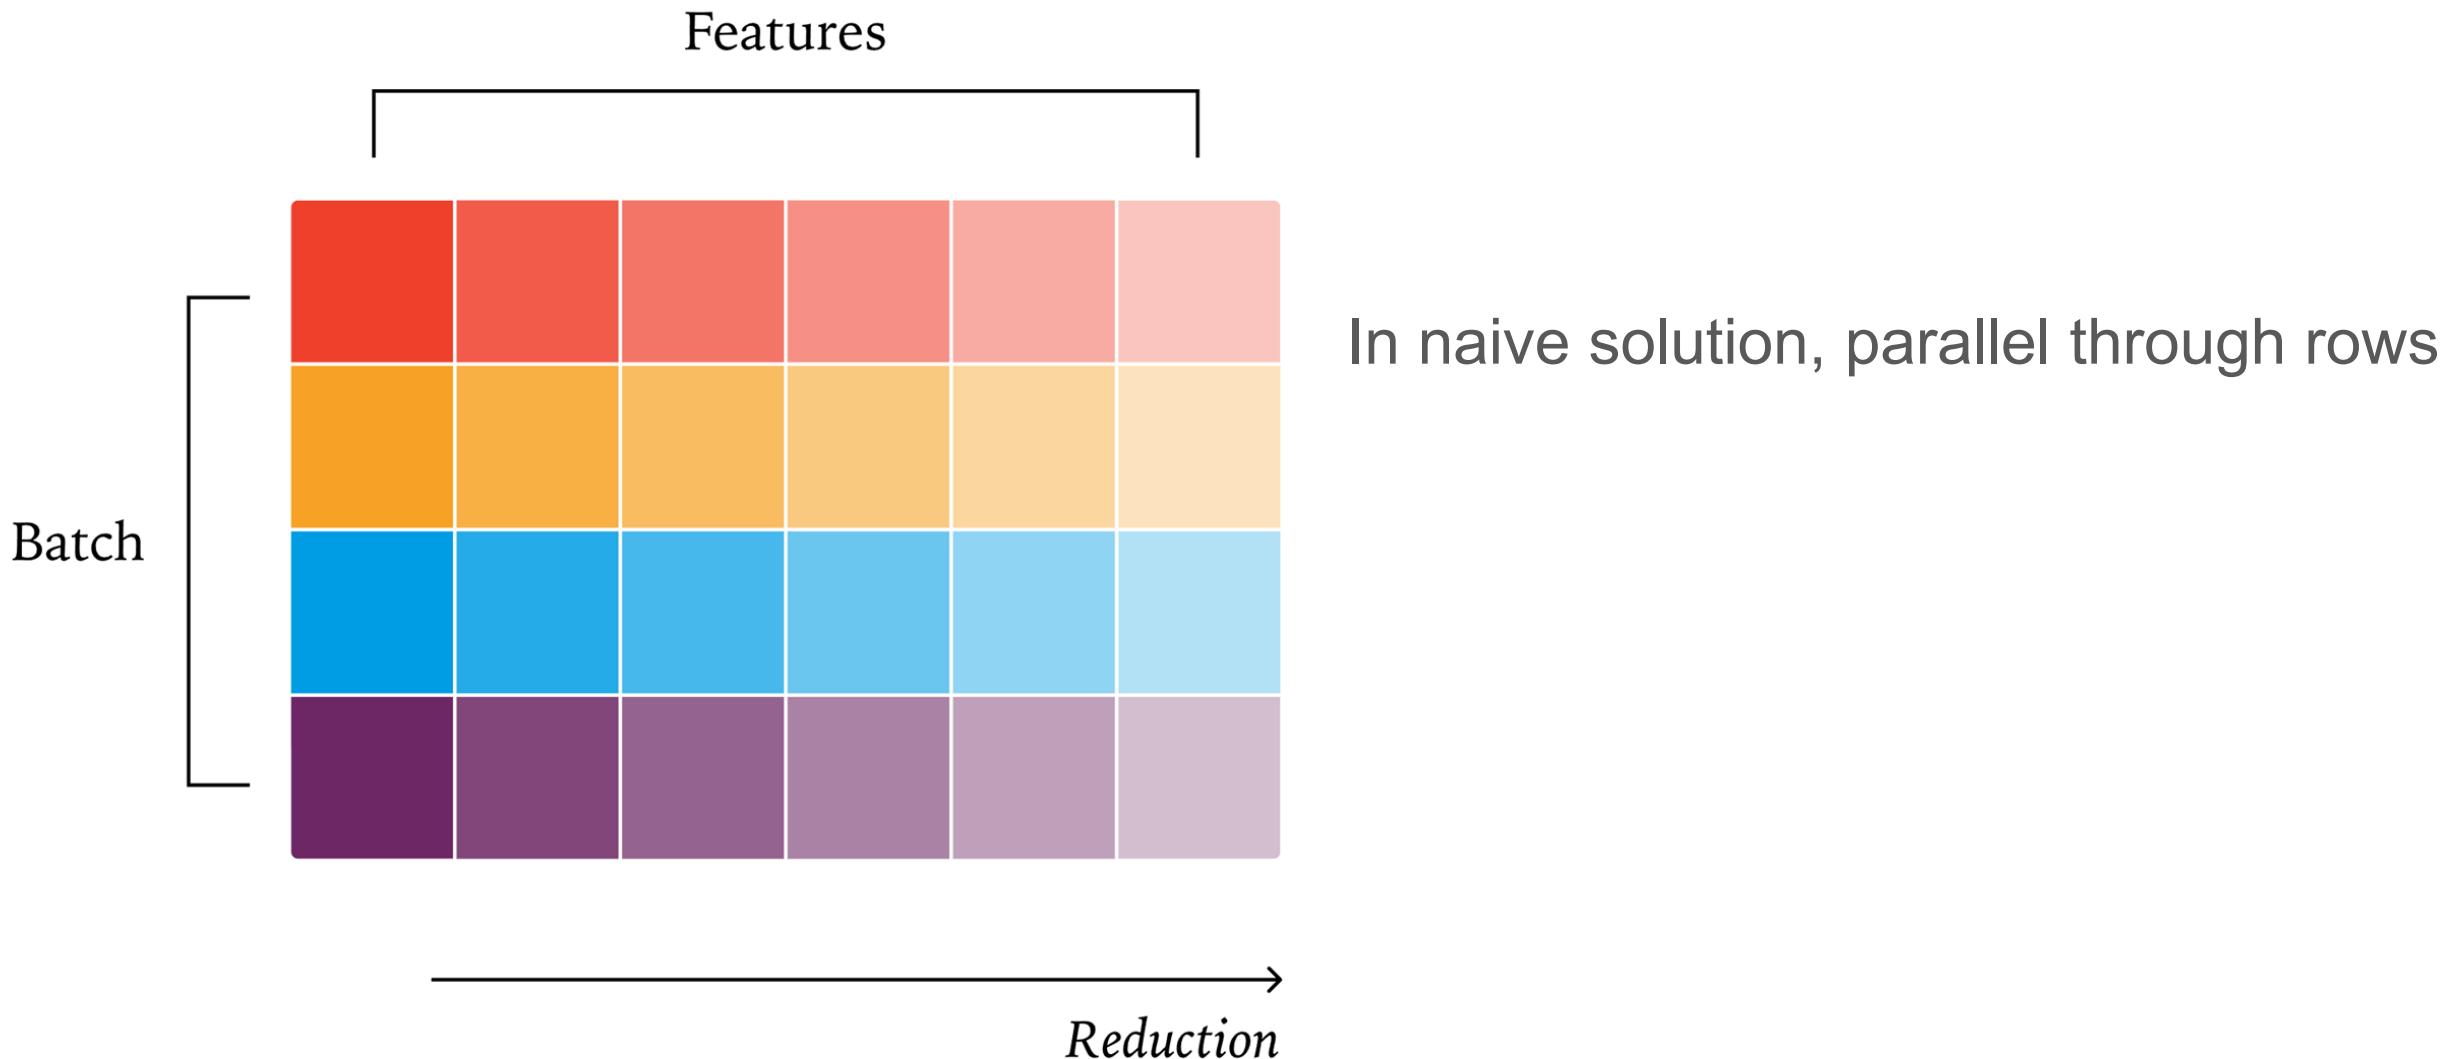
Make every kernel involving reductions batch invariant

For LLM inference, **three kernels dominate**

- RMSNorm
- Matrix Mul
- Attention

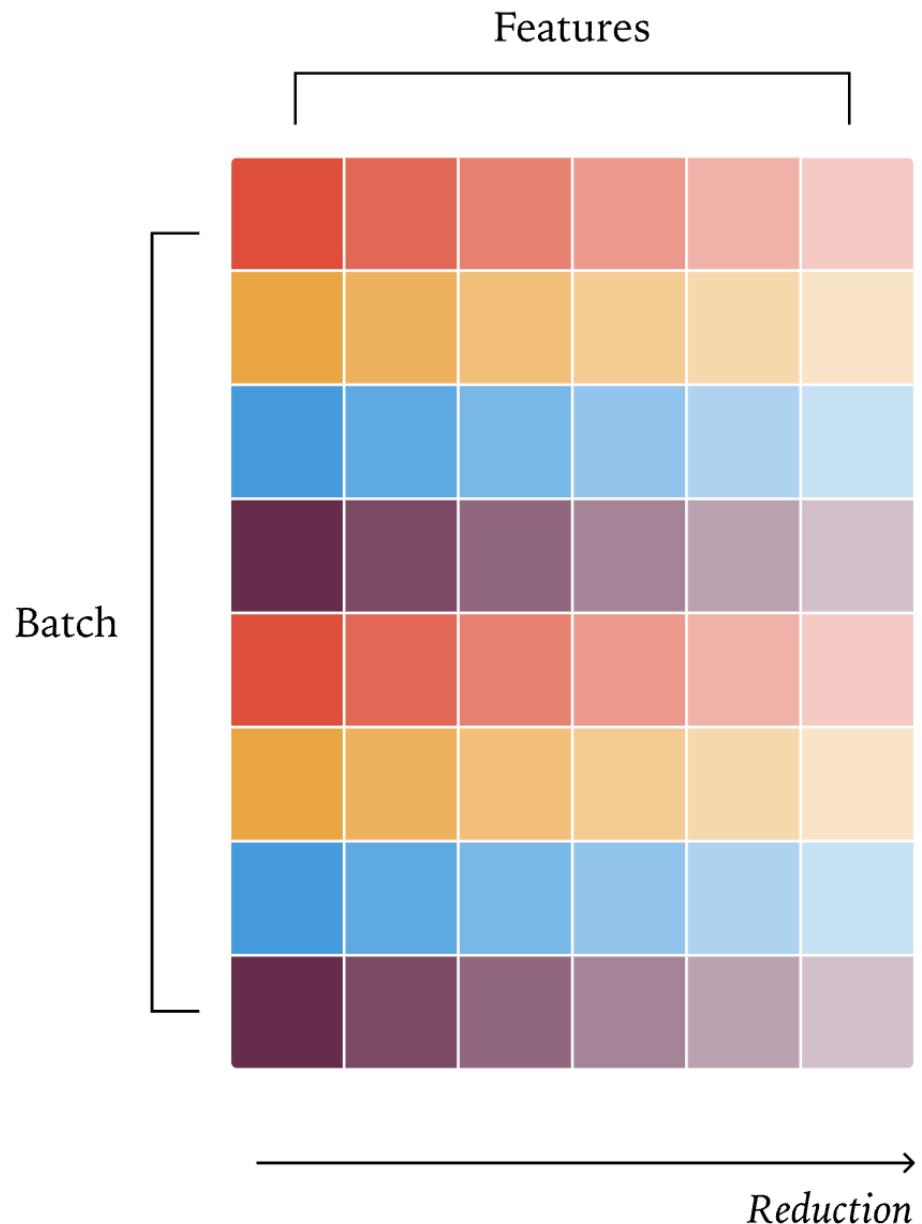
## Key 1: RMSNorm

$$y = x \cdot \text{rsqrt} \left( \text{mean}(x^2) + \epsilon \right) \cdot w$$



## Key 1: RMSNorm

$$y = x \cdot \text{rsqrt} \left( \text{mean}(x^2) + \epsilon \right) \cdot w$$

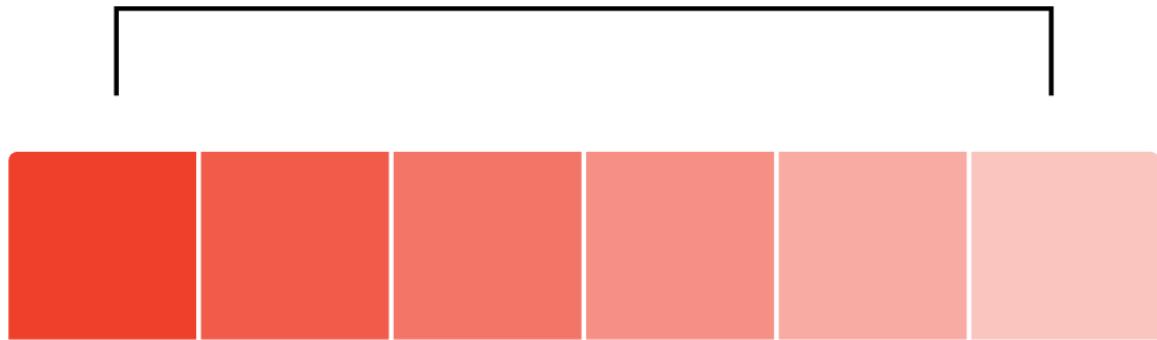


In large batch, this works well

## Key 1: RMSNorm

$$y = x \cdot \text{rsqrt} \left( \text{mean}(x^2) + \epsilon \right) \cdot w$$

Features



**In small batch, wasting  
compute resources**

## Key 1: RMSNorm

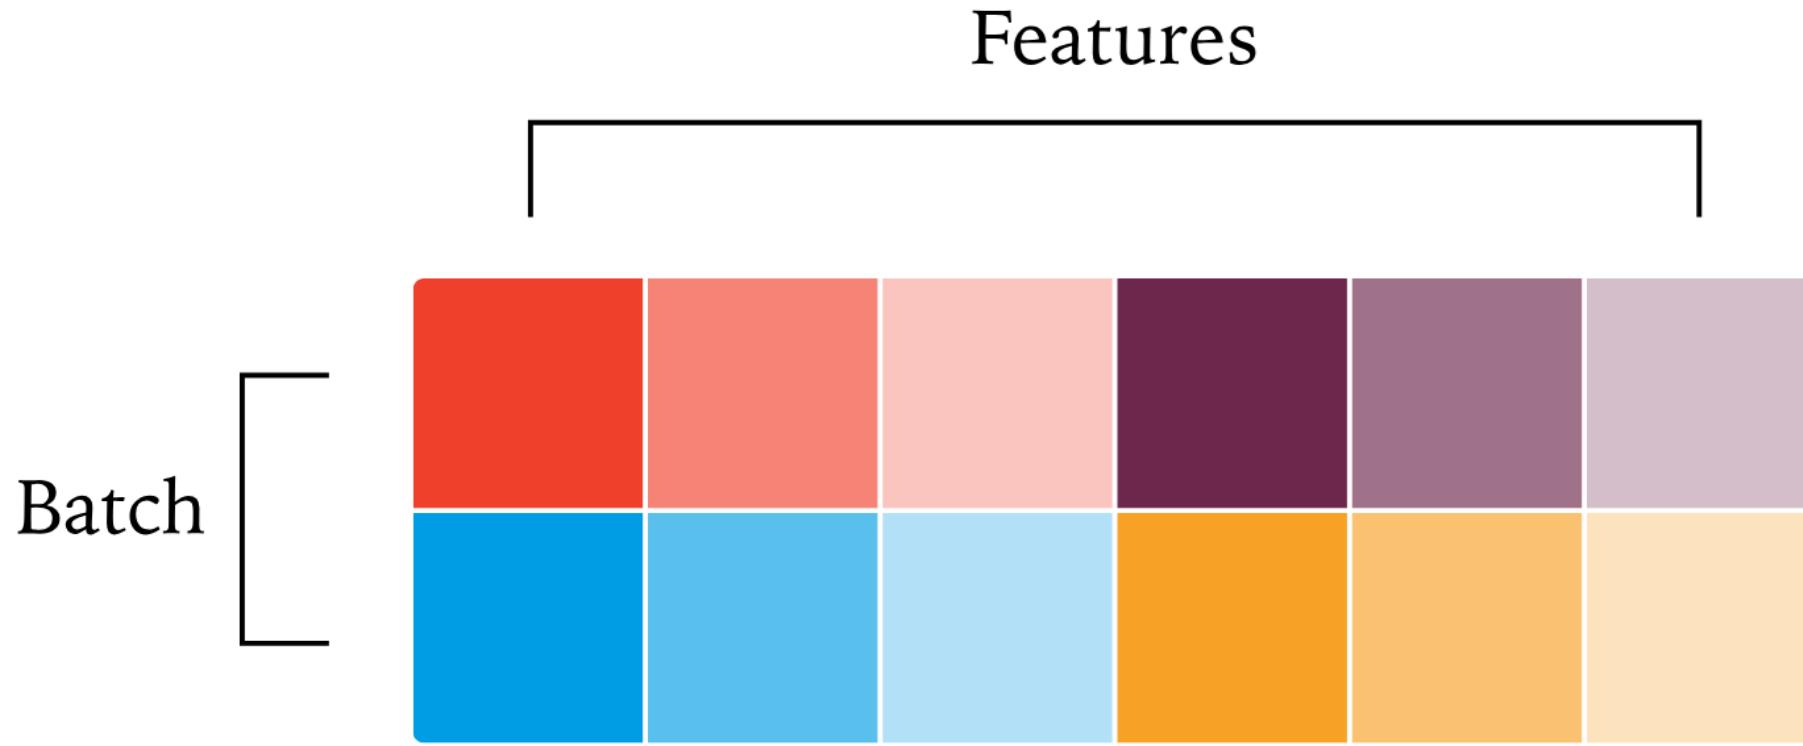
$$y = x \cdot \text{rsqrt} \left( \text{mean}(x^2) + \epsilon \right) \cdot w$$



In **small batch**, we usually optimize using Split-K

## Key 1: RMSNorm

$$y = x \cdot \text{rsqrt} \left( \text{mean}(x^2) + \epsilon \right) \cdot w$$

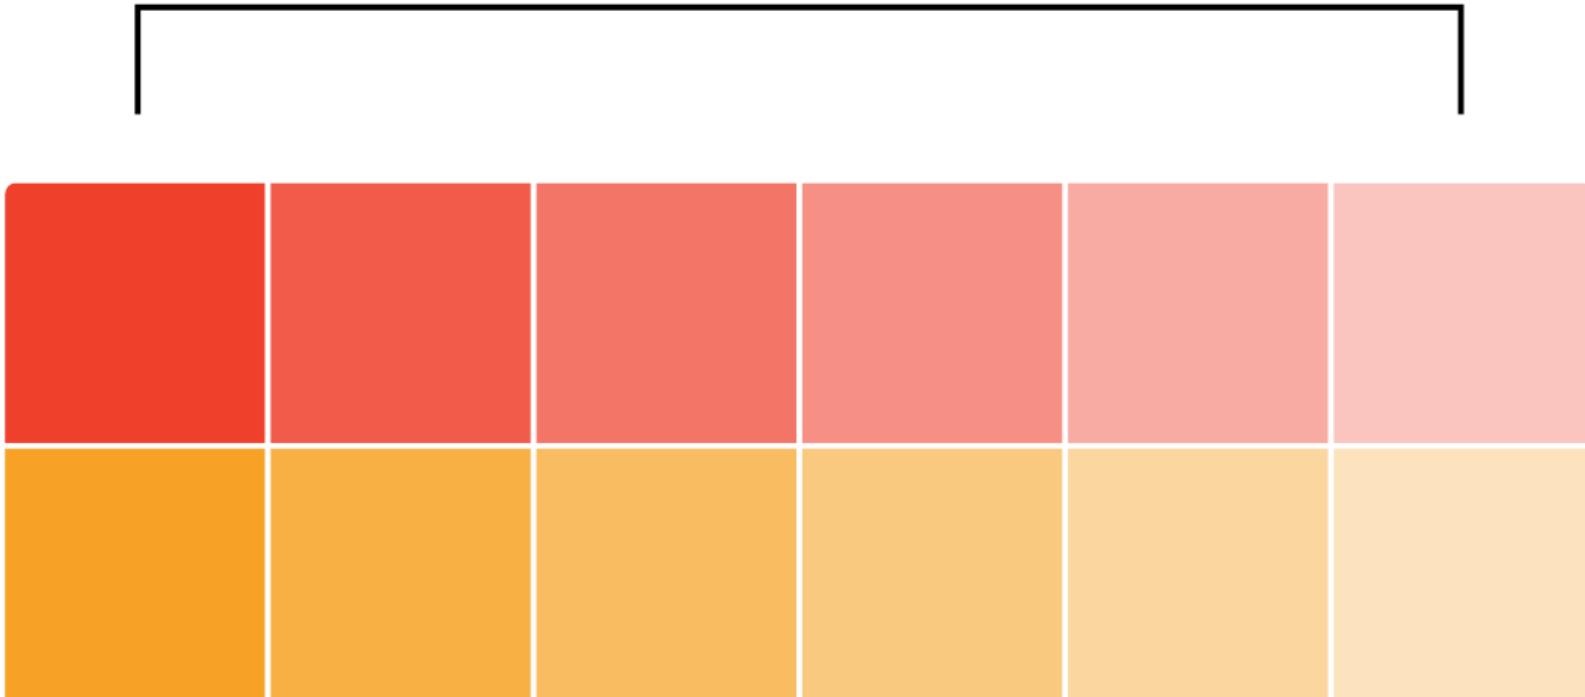


However, as  $x$  grows, the reduction order changes

## Key 1: RMSNorm

$$y = x \cdot \text{rsqrt} \left( \text{mean}(x^2) + \epsilon \right) \cdot w$$

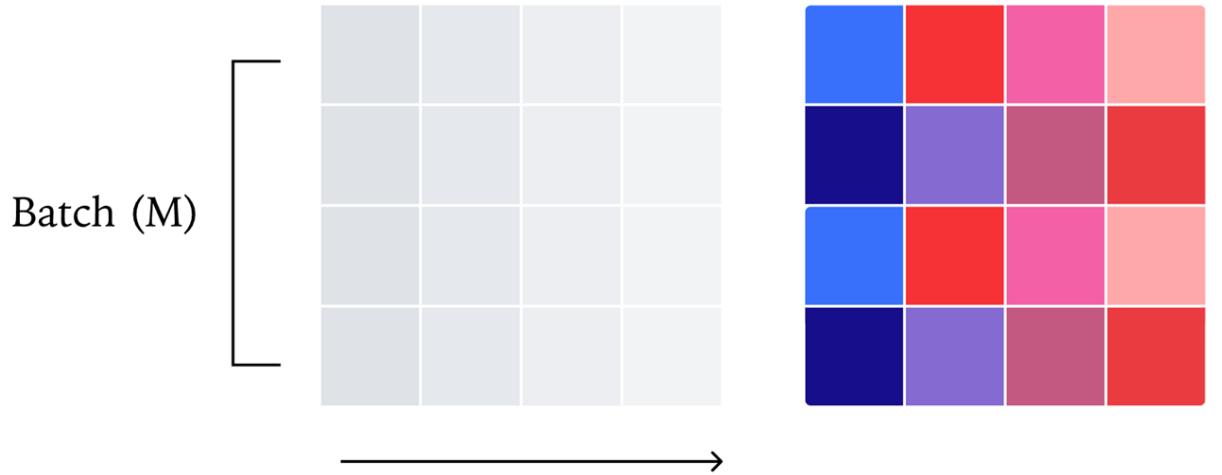
Features



The easiest way is just not to do the split and keep the order

## Key 2: Matrix Mul

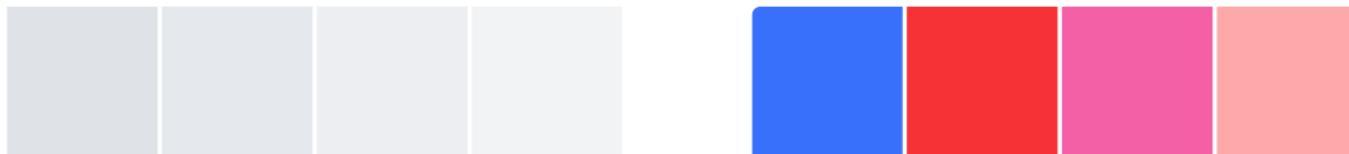
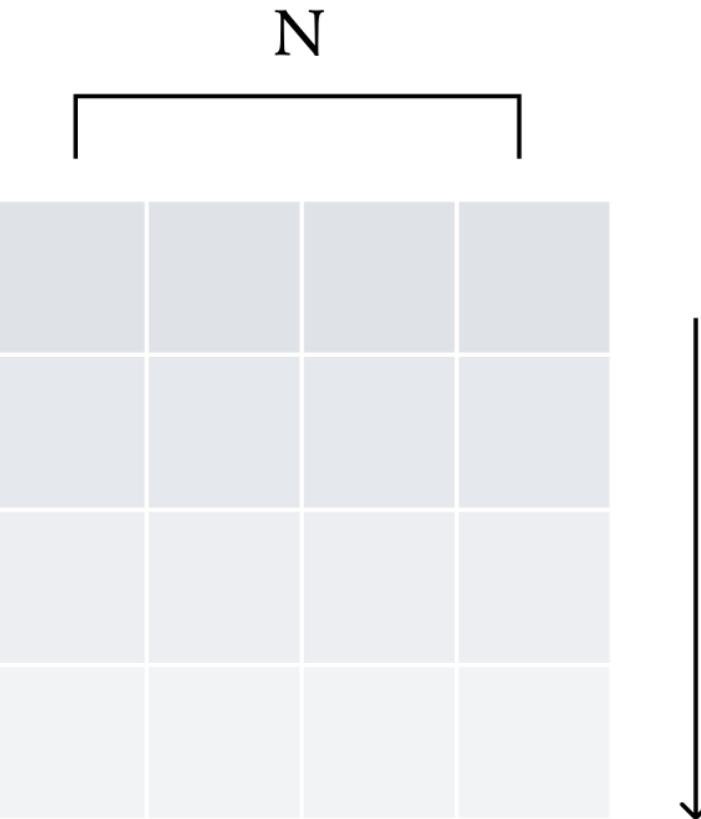
$$Y = XW, \quad X \in \mathbb{R}^{M \times K}, W \in \mathbb{R}^{K \times N}, Y \in \mathbb{R}^{M \times N}$$



The naive matrix mul follows the batch invariant logic, good for large batch

## Key 2: Matrix Mul

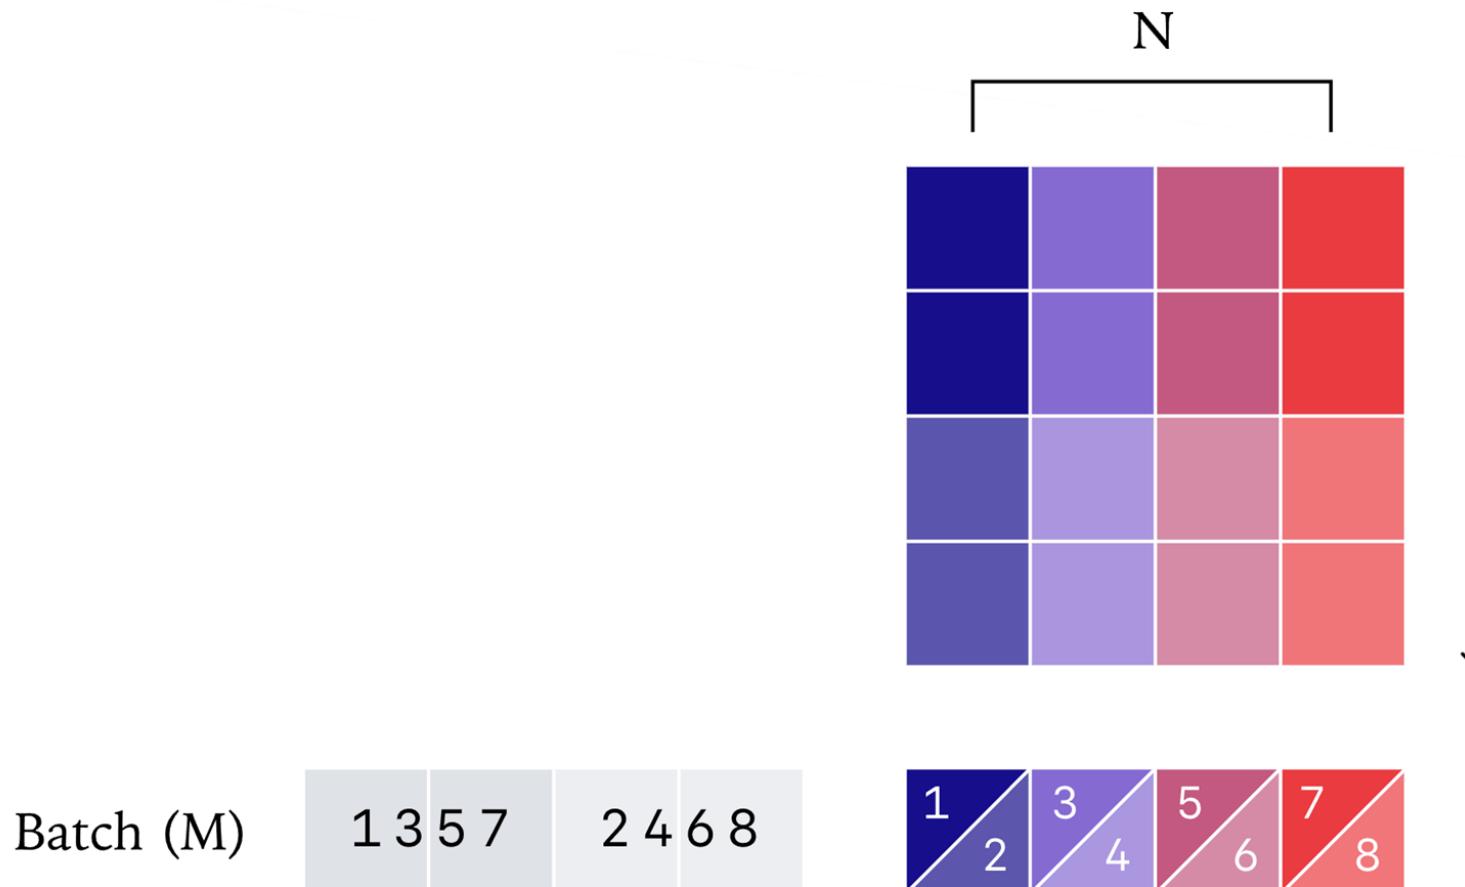
$$Y = XW, \quad X \in \mathbb{R}^{M \times K}, \quad W \in \mathbb{R}^{K \times N}, \quad Y \in \mathbb{R}^{M \times N}$$



Similar problem when batch size is small

## Key 2: Matrix Mul

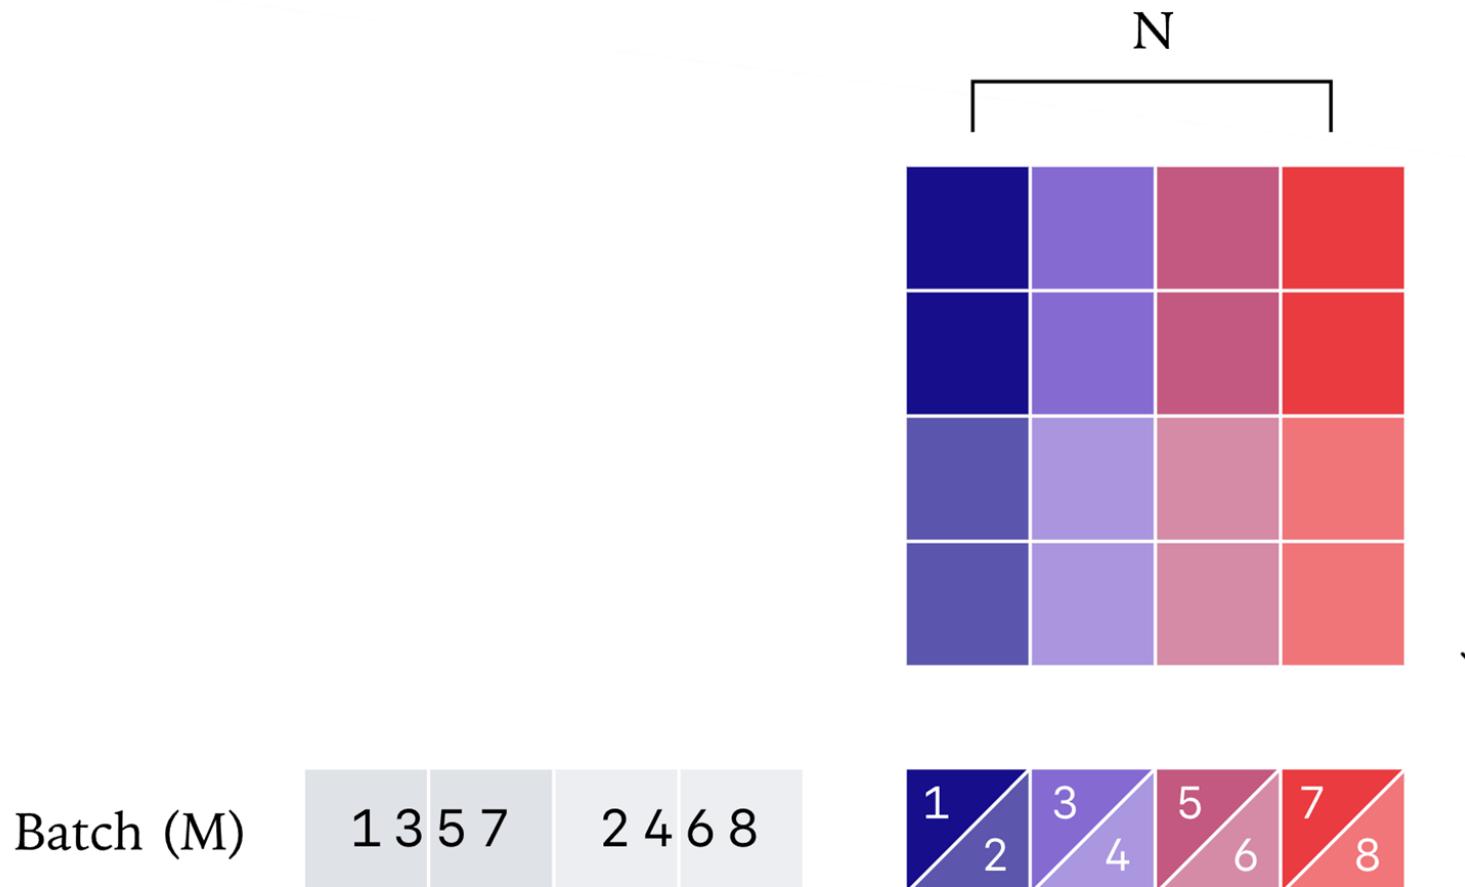
$$Y = XW, \quad X \in \mathbb{R}^{M \times K}, \quad W \in \mathbb{R}^{K \times N}, \quad Y \in \mathbb{R}^{M \times N}$$



In **small batch**, we usually optimize using Split-K

## Key 2: Matrix Mul

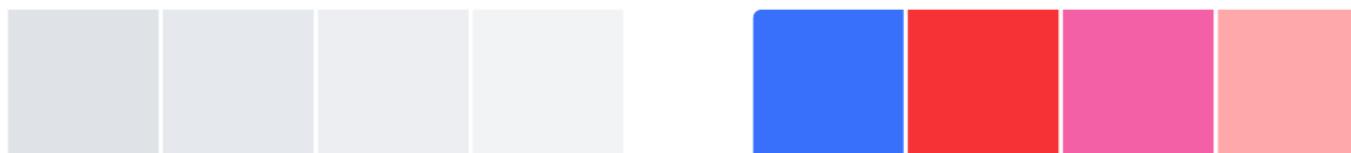
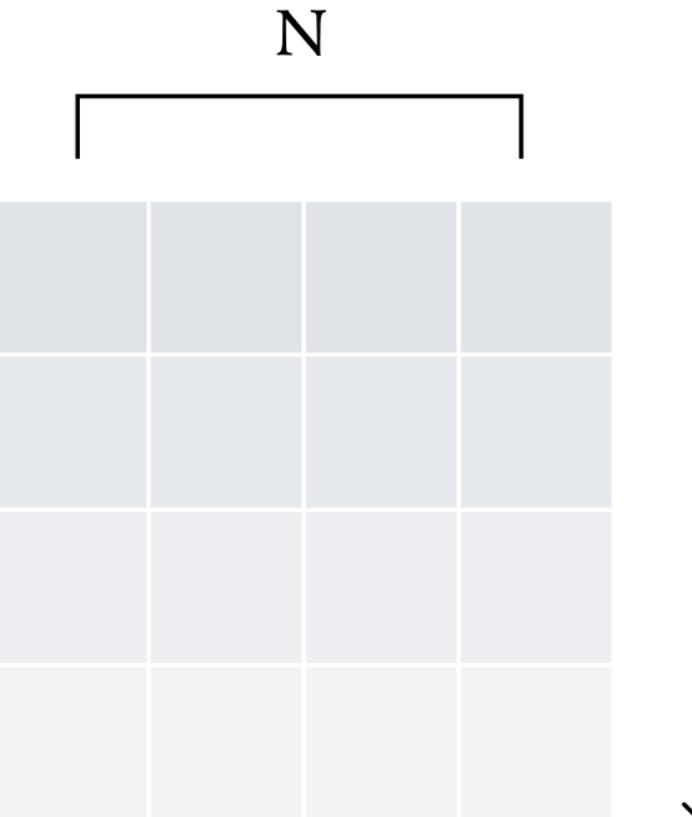
$$Y = XW, \quad X \in \mathbb{R}^{M \times K}, \quad W \in \mathbb{R}^{K \times N}, \quad Y \in \mathbb{R}^{M \times N}$$



This breaks the **order of reduction** as  $K$  grows

## Key 2: Matrix Mul

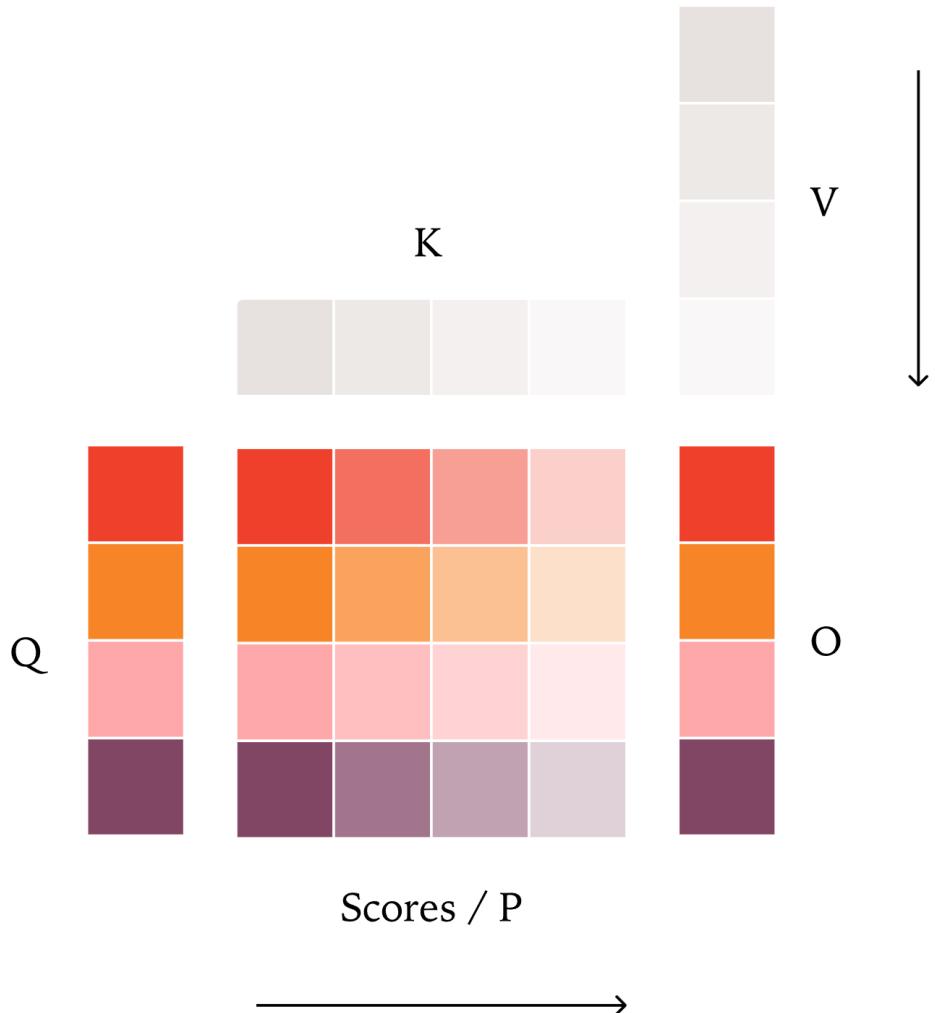
$$Y = XW, \quad X \in \mathbb{R}^{M \times K}, \quad W \in \mathbb{R}^{K \times N}, \quad Y \in \mathbb{R}^{M \times N}$$



Solution: compile one kernel configuration and use that for all shapes (fixed tile, fixed split-k etc)

## Key 3: Attention

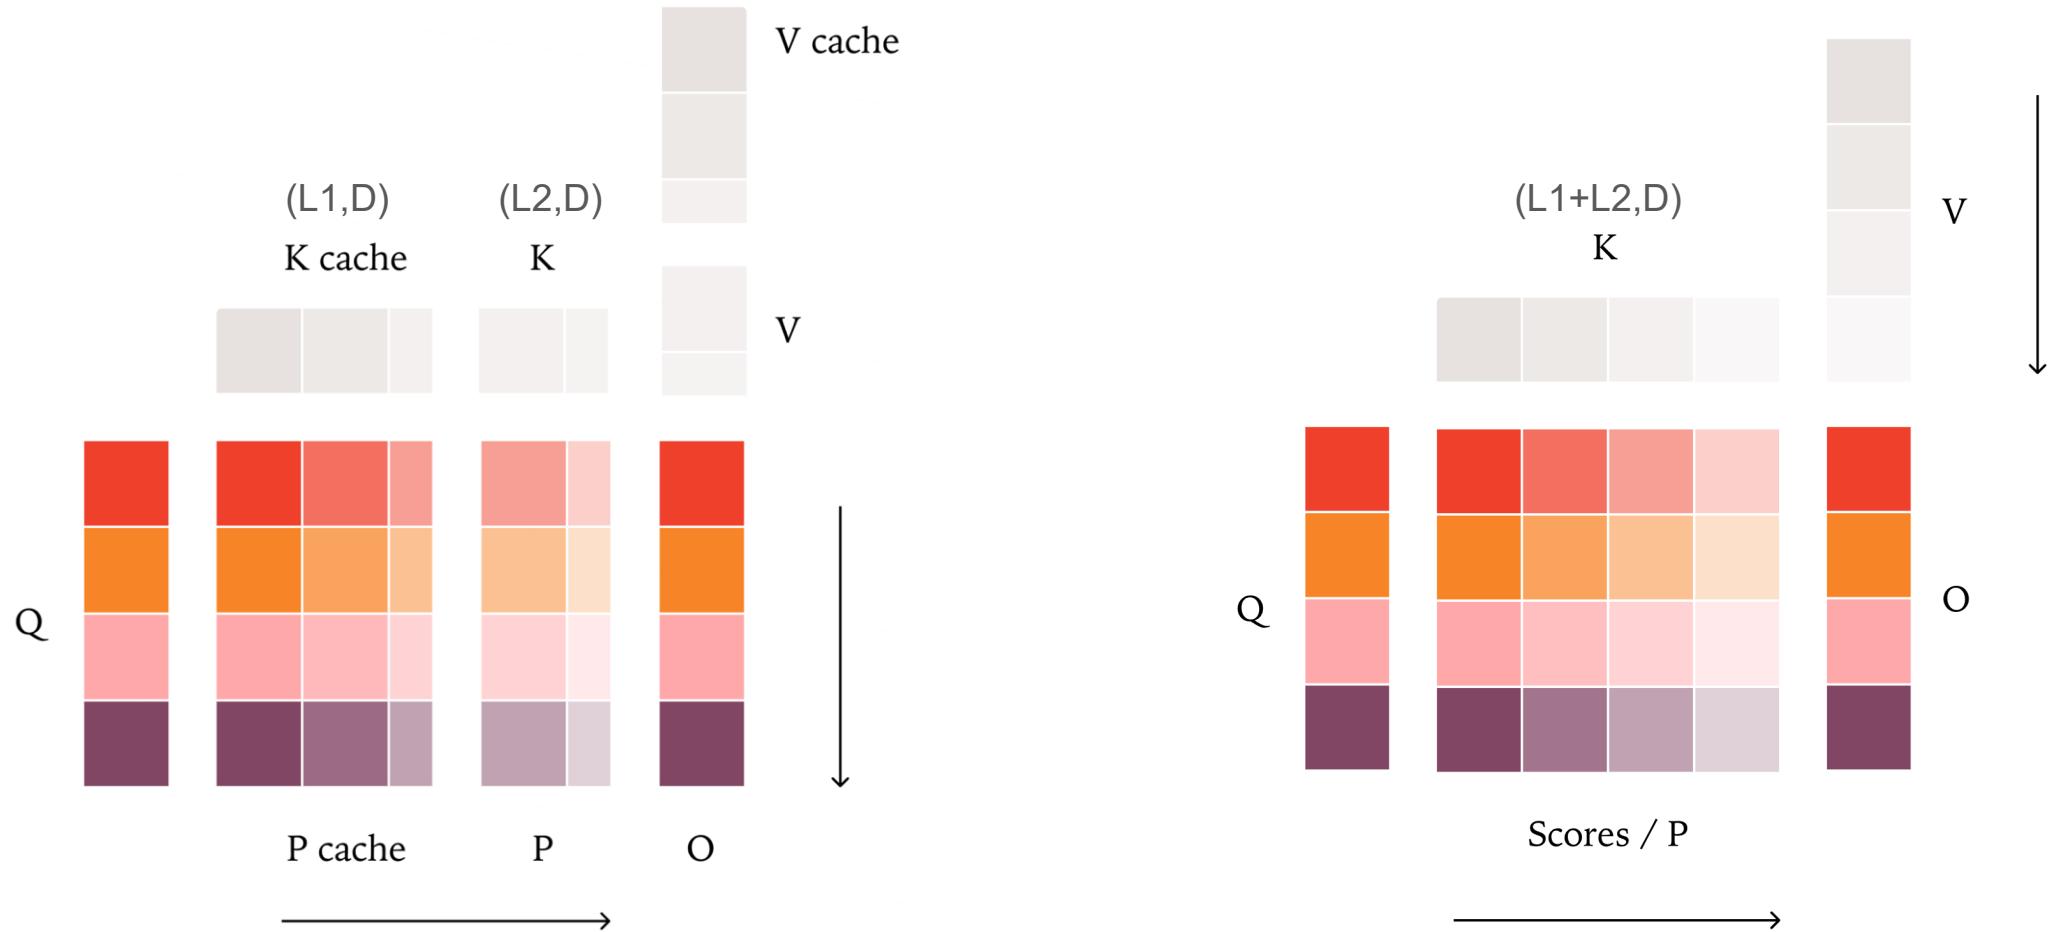
$$O = \text{softmax}\left(\frac{QK^\top}{\sqrt{d}}\right)V$$



Parallel along  $Q$  in naive solution,  
which is good for batch invariance

## Key 3: Attention

$$O = \text{softmax}\left(\frac{QK^\top}{\sqrt{d}}\right)V$$

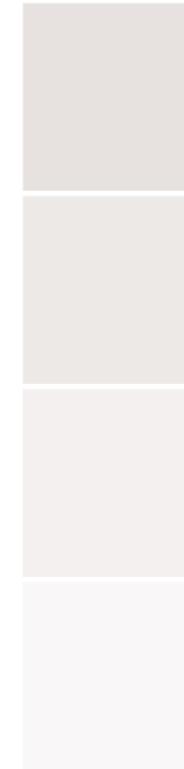


Calculation for  $K$ , then added to the  $K-V$  cache

## Key 3: Attention

$$O = \text{softmax}\left(\frac{QK^\top}{\sqrt{d}}\right)V$$

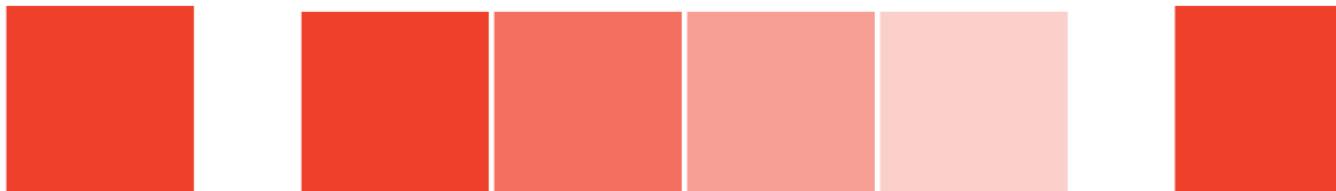
K



V



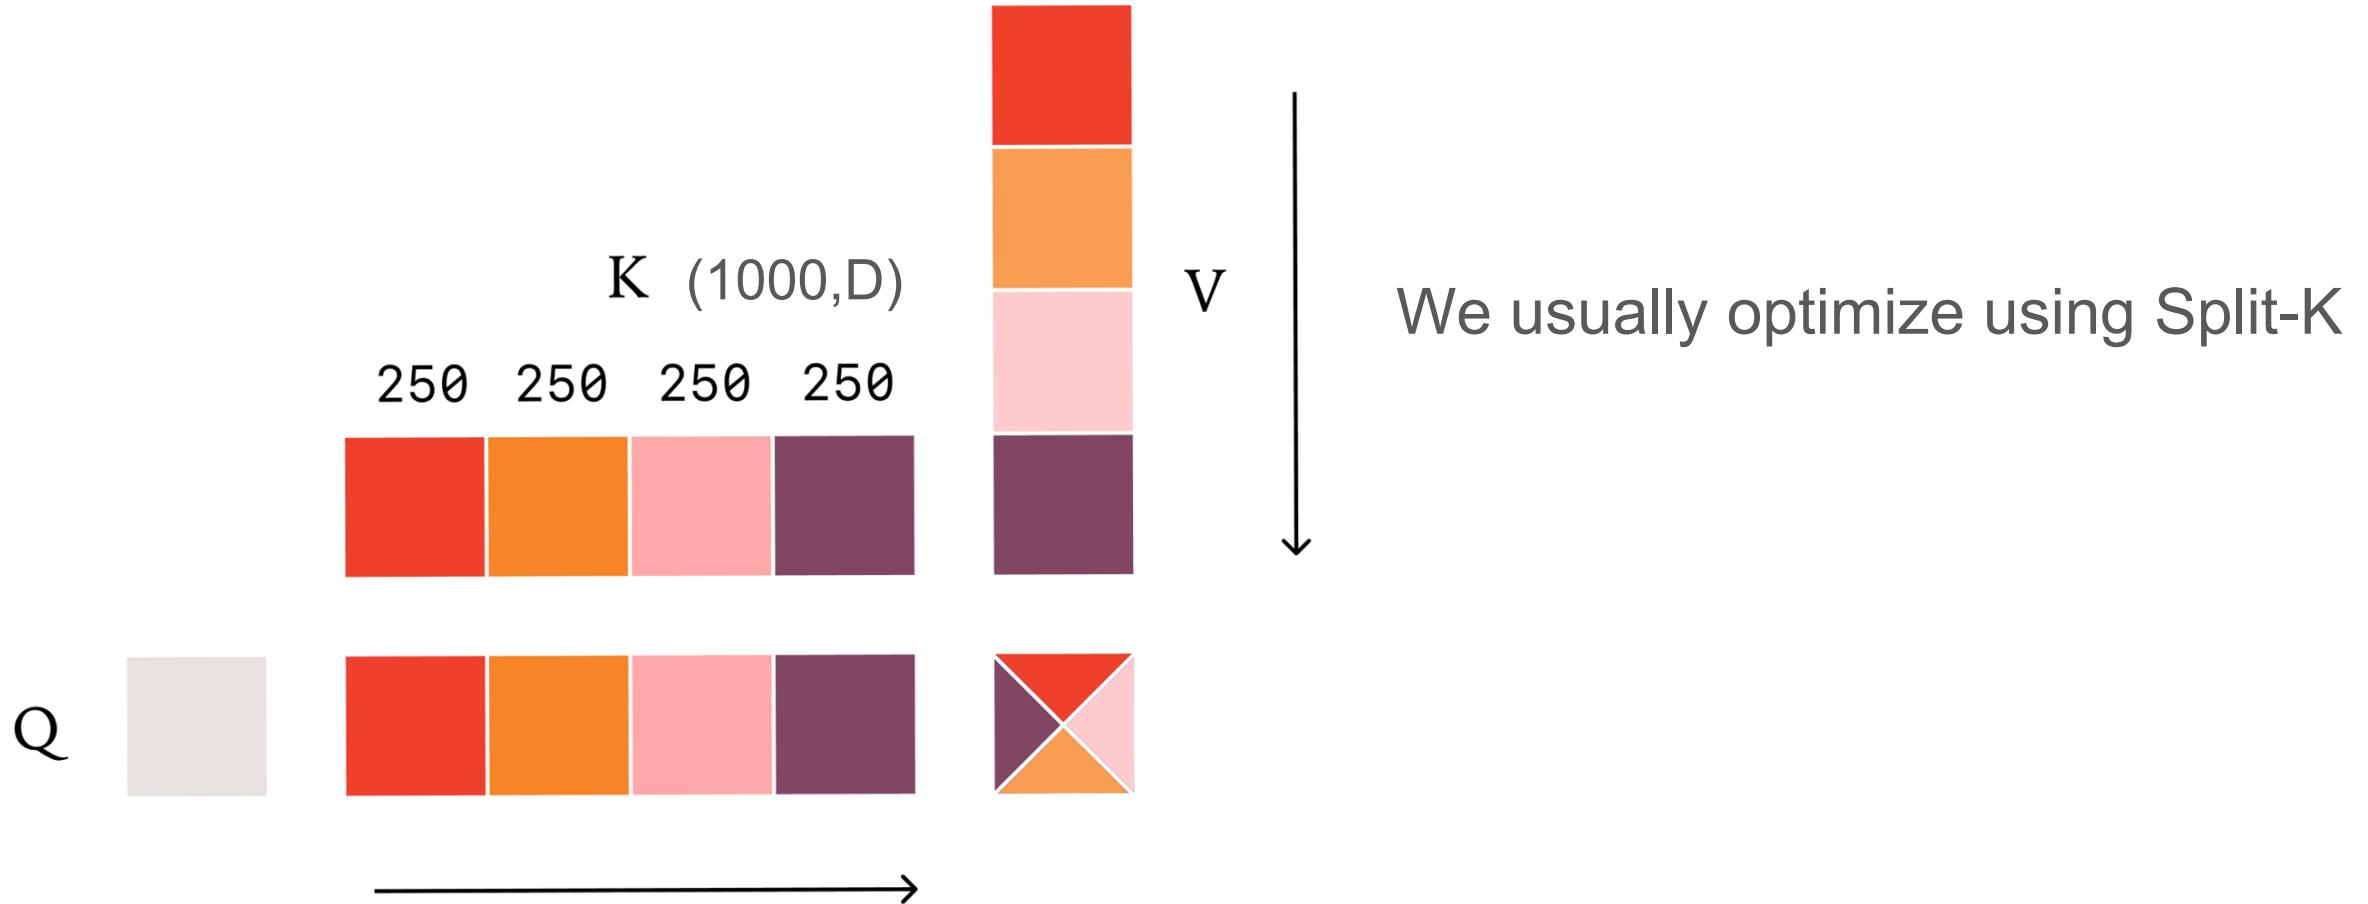
Q



Similar problem when batch is small

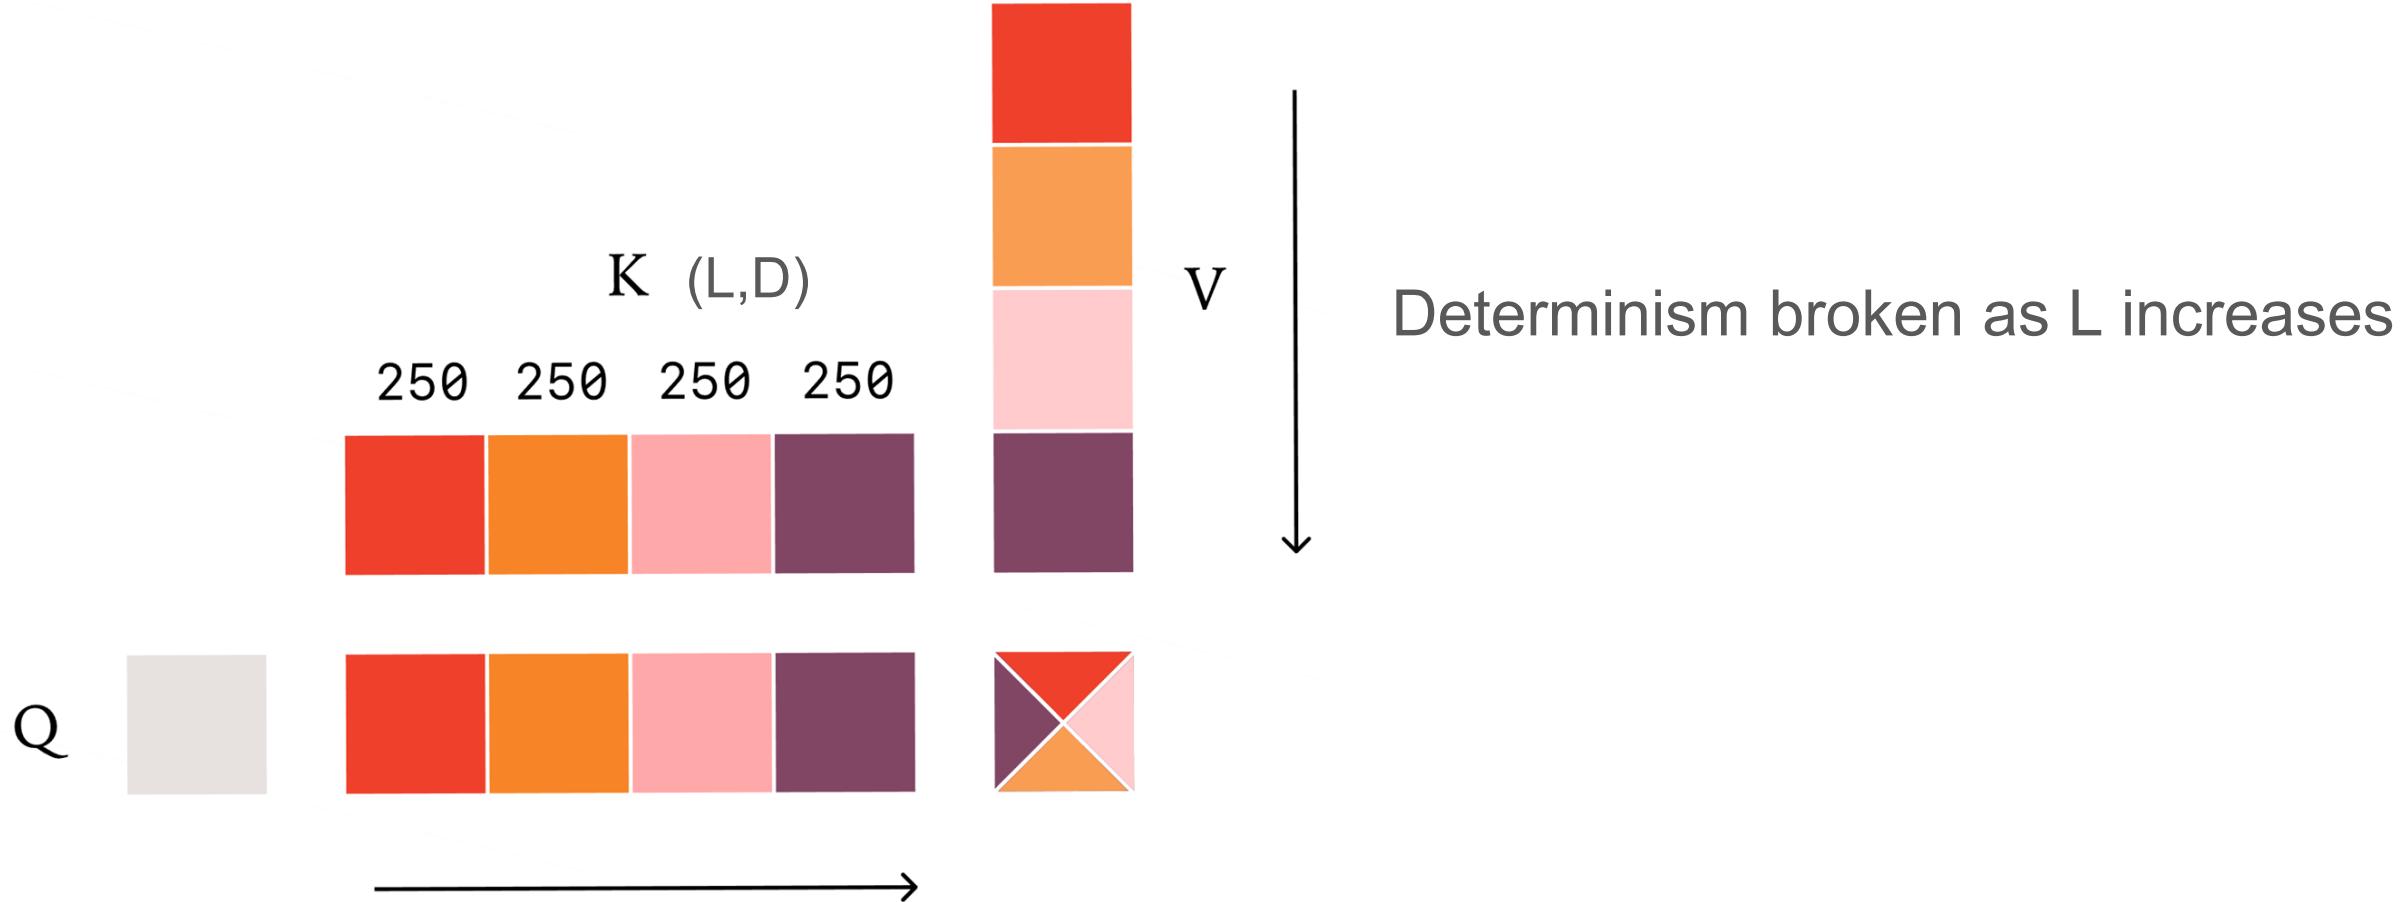
## Key 3: Attention

$$O = \text{softmax}\left(\frac{QK^\top}{\sqrt{d}}\right)V$$



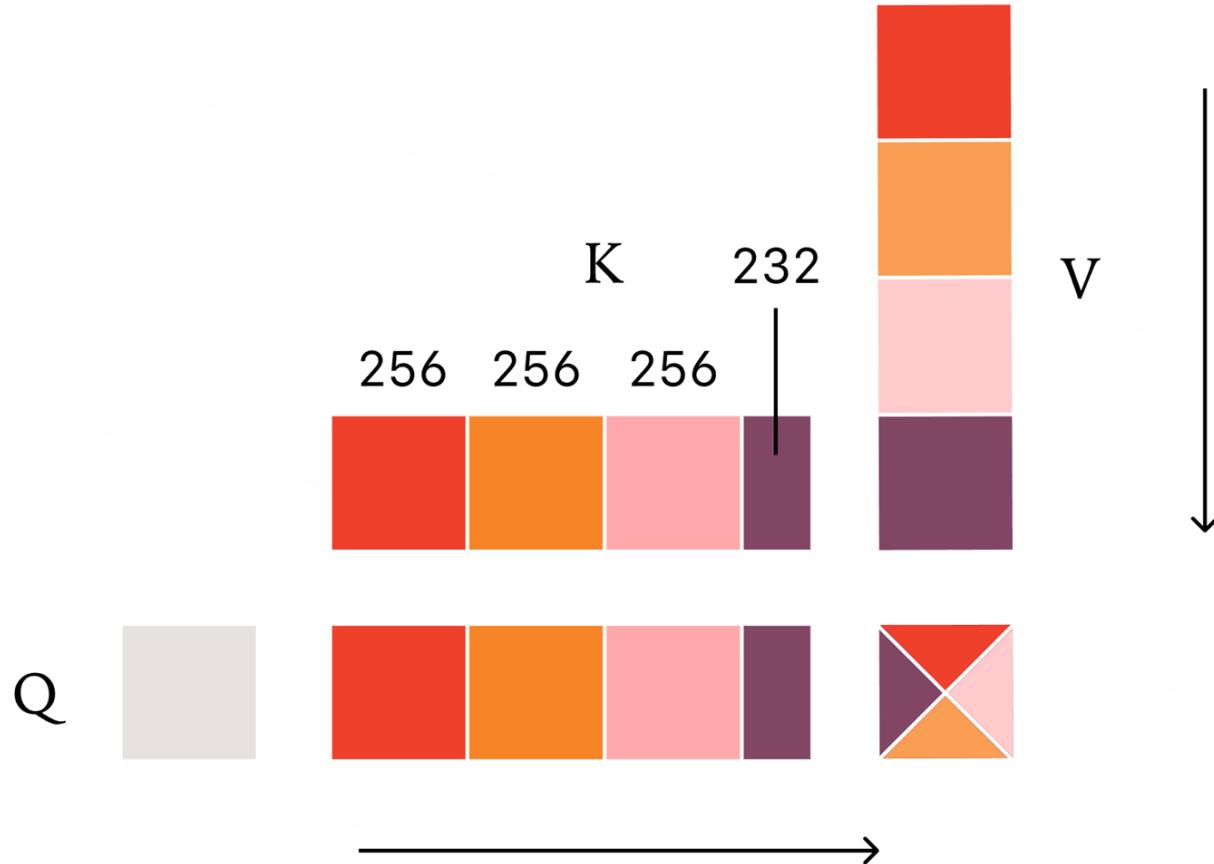
### Key 3: Attention

$$O = \text{softmax}\left(\frac{QK^\top}{\sqrt{d}}\right)V$$



## Key 3: Attention

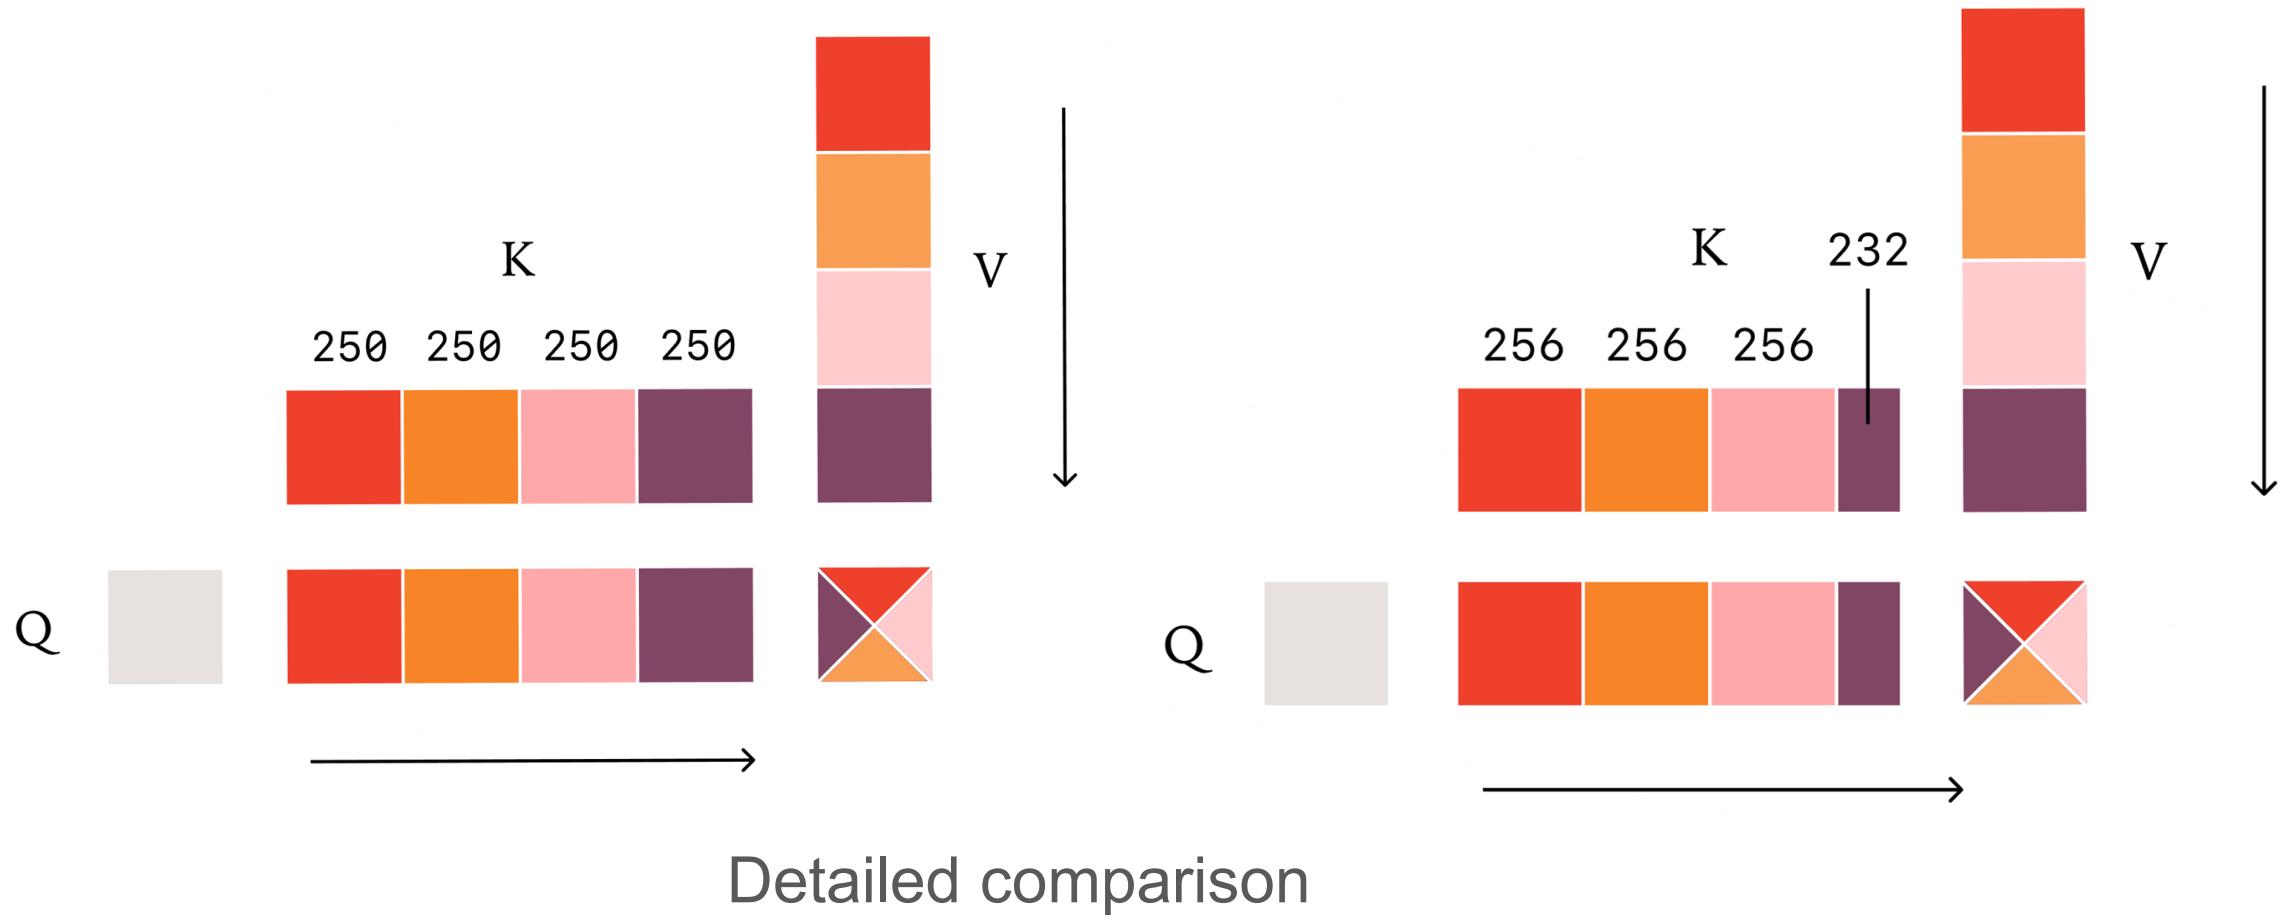
$$O = \text{softmax}\left(\frac{QK^\top}{\sqrt{d}}\right)V$$



So we should use a fixed split size  
e.g., 256

## Key 3: Attention

$$O = \text{softmax}\left(\frac{QK^\top}{\sqrt{d}}\right)V$$



## Batch invariance summary

```
$ VLLM_BATCH_INVARIANT=0 python invariance_test.py  
Logprob difference: 0.0000052527  
$ VLLM_BATCH_INVARIANT=1 python invariance_test.py  
Logprob difference: 0.0000000000
```

Every kernel is batch invariant, with fixed reduction order

\*\*Simplified due to time constraints.

# Tracking the Latest Updates

The screenshot shows a GitHub issue page for a project named "vllm-project / vllm". The issue is titled "[Feature]: Batch Invariant Feature and Performance Optimization #27433". The issue was opened by "yewentao256" on Oct 23, 2025, and last edited by "yewentao256". The status is "Open".

**Assignees:** PaulZhang12, bwasti, yewentao256

**Labels:** feature request

**Type:** No type

**Projects:** Batch-invariant Inference

**Status:** In Progress

**Milestone:** No milestone

**Relationships:** None

**Issue Content:**

**Summary:** The feature, motivation and pitch

We have basically support Batch Invariant based on <https://thinkingmachines.ai/blog/defeating-nondeterminism-in-llm-inference/>

**Batch-invariant Inference (view)**

But there are still some work to be done, so here is the issue to track the work

**TODOS:**

- Basic framework ↗ Kernel-override Determinism [1/n] #25603 @bwasti
- FlashInfer support ↗ [unrevert] Add batch invariant kernel override for FlashInfer backend [2/n] #26373 @bwasti
- Deepseek-v3 ↗ Deepseek-v3 Batch Invariant on 8xH100 #26609 @bwasti
- DeepGEMM on Blackwell ↗ [Feature] Batch Invariant: Support DeepGEMM and Blackwell #27127 @yewentao256
- Batch Invariant for R1 TP 8 on Blackwell ↗ [Feature] Batch Invariant for R1 TP 8 on Blackwell #27229 @yewentao256
- Torch compile & Cuda Graph support ↗ [Feature] Batch invariant torch.compile #27660 @PaulZhang12
- Usability & Documentation @bwasti ↗ Batch invariance doc #27839

<https://github.com/vllm-project/vllm/issues/27433>

Features, optimizations, model validations...  
Any help is greatly appreciated!

# Contribution to Batch Invariance in vLLM

## Tested Models

Batch invariance has been tested and verified on the following models:

- **DeepSeek series:** deepseek-ai/DeepSeek-V3 , deepseek-ai/DeepSeek-V3-0324 , deepseek-ai/DeepSeek-R1 , deepseek-ai/DeepSeek-V3.1
- **Qwen3 (Dense):** Qwen/Qwen3-1.7B , Qwen/Qwen3-8B
- **Qwen3 (MoE):** Qwen/Qwen3-30B-A3B , Qwen/Qwen3-Next-80B-A3B-Instruct
- **Llama 3:** meta-llama/Llama-3.1-8B-Instruct , meta-llama/Llama-3.2-1B-Instruct



Help needed for validations of more models.

1. Test a model using the current script
2. Submit a PR updating the document

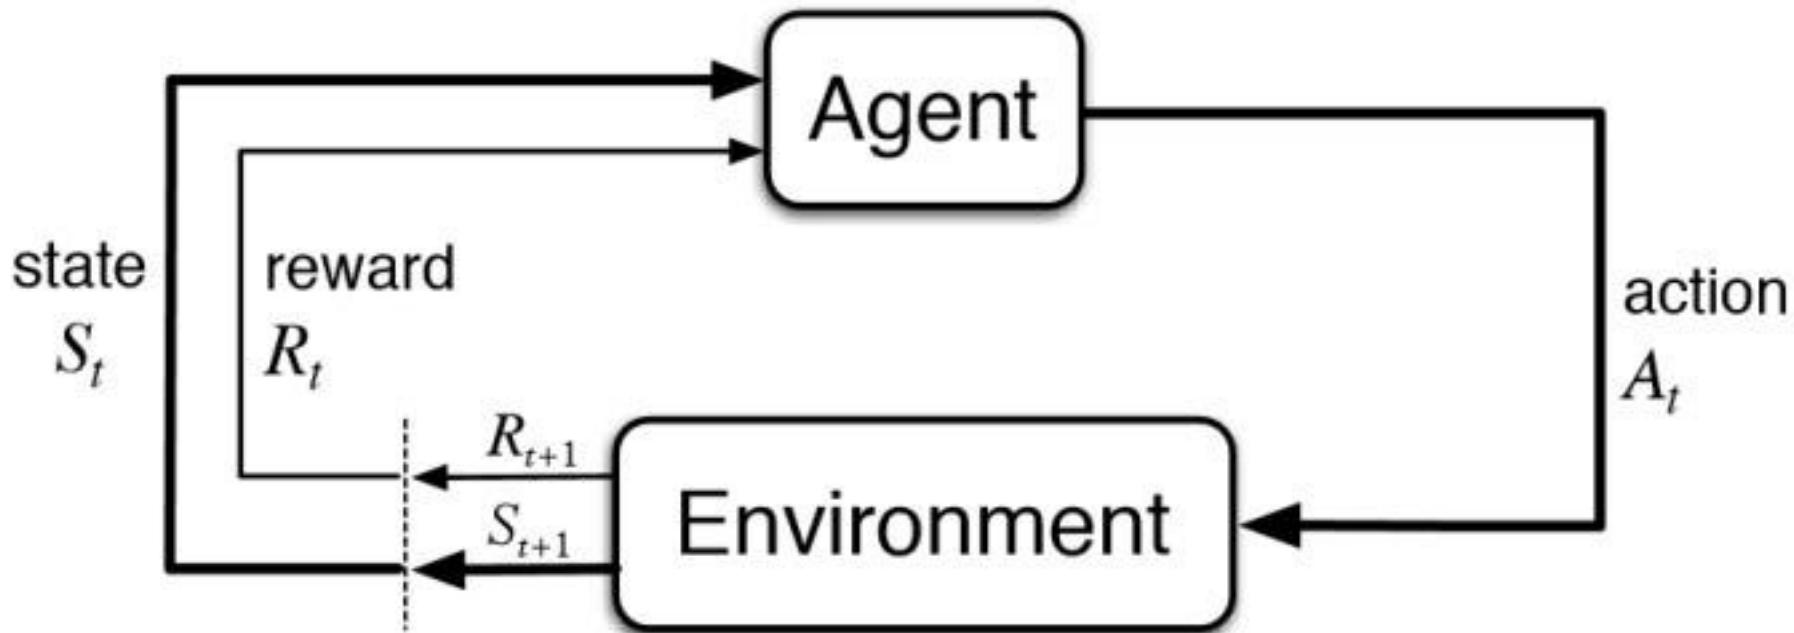
<https://github.com/vllm-project/vllm/issues/27433>

## References

- [1]: [https://docs.vllm.ai/en/latest/features/batch\\_invariance/](https://docs.vllm.ai/en/latest/features/batch_invariance/) "Batch Invariance - vLLM"
- [2]: <https://vllm.ai/> "vLLM"
- [3]: <https://arxiv.org/html/2506.09501v2> "Understanding and Mitigating Numerical Sources of Nondeterminism in LLM Inference"
- [4]: <https://thinkingmachines.ai/blog/defeating-nondeterminism-in-lm-inference/> "Defeating Nondeterminism in LLM Inference - Thinking Machines Lab"
- [5]: <https://github.com/vllm-project/vllm/issues/27433> "[Feature]: Batch Invariant Feature and Performance Optimization · Issue #27433 · vllm-project/vllm · GitHub"
- [6]: <https://blog.vllm.ai/2025/11/10/bitwise-consistent-train-inference.html> "No More Train-Inference Mismatch: Bitwise Consistent On-Policy Reinforcement Learning with vLLM and TorchTitan | vLLM Blog"
- [7]: <https://insujang.github.io/2024-01-07/lm-inference-continuous-batching-and-pagedattention/> LLM Inference: Continuous Batching and PagedAttention

# **Bitwise Consistent On-Policy Reinforcement Learning with vLLM + TorchTitan**

# Reinforcement Learning (Basics)



Ignore all these variable letter choices :P

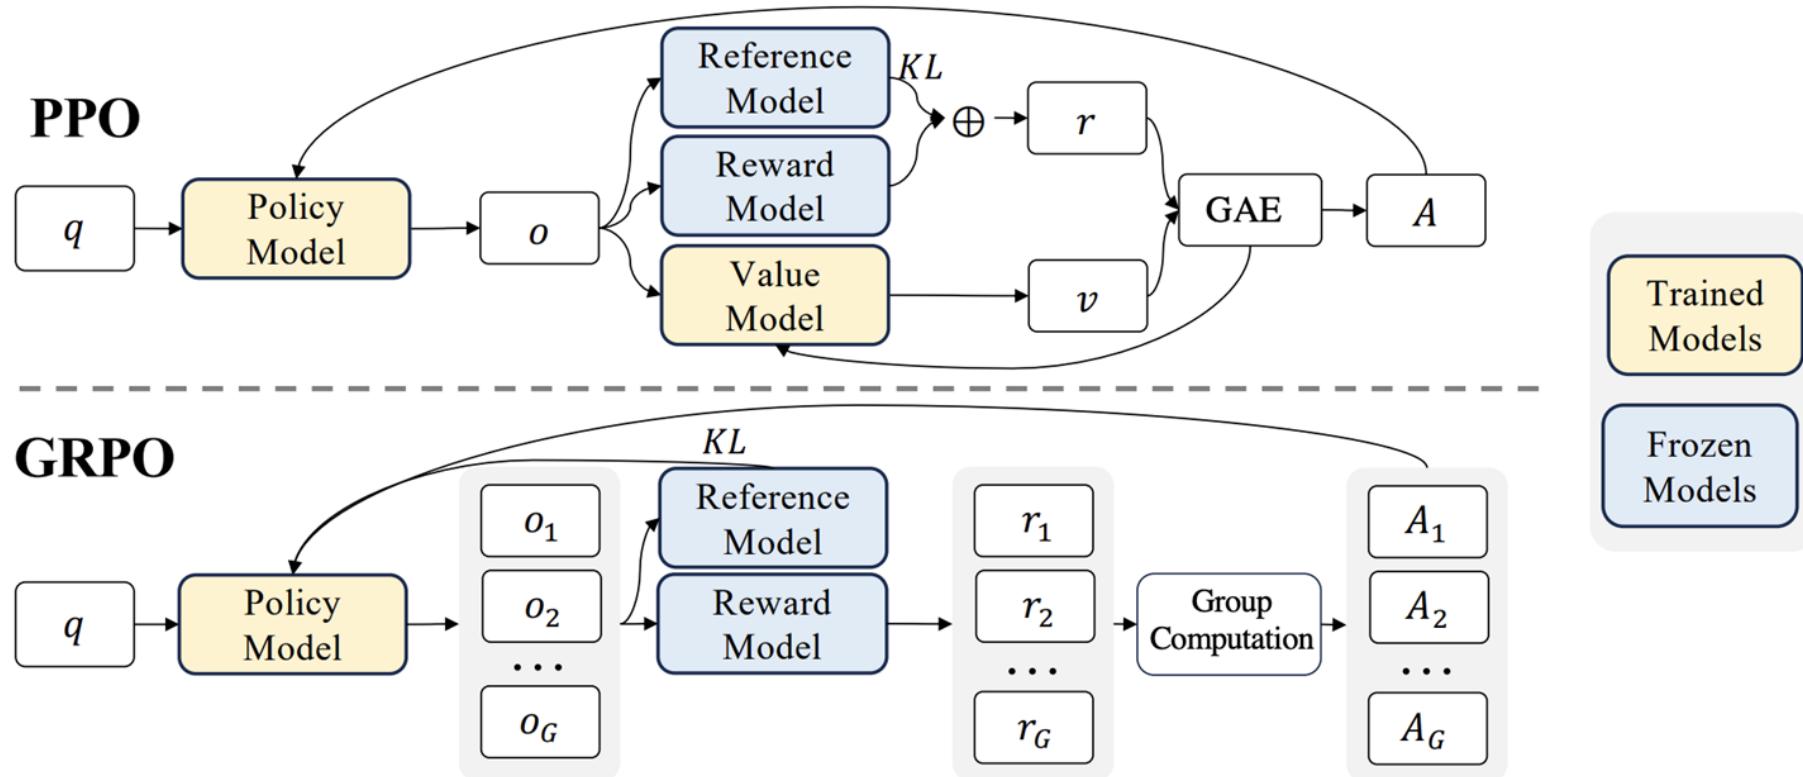
## Reinforcement Learning (PPO)

$$\hat{g} = \hat{\mathbb{E}}_t \left[ \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \hat{A}_t \right]$$

<https://arxiv.org/abs/1707.06347>

Advantage calculation is tricky

# Reinforcement Learning (PPO -> GRPO)



<https://arxiv.org/pdf/2402.03300>

# Reinforcement Learning (IS Weights)

Running off-policy needs a correction term!

$$\begin{aligned} J(\theta) &= \sum_{s \in \mathcal{S}} \rho^{\pi_{\theta_{\text{old}}}} \sum_{a \in \mathcal{A}} (\pi_{\theta}(a|s) \hat{A}_{\theta_{\text{old}}}(s, a)) \\ &= \sum_{s \in \mathcal{S}} \rho^{\pi_{\theta_{\text{old}}}} \sum_{a \in \mathcal{A}} \left( \beta(a|s) \frac{\pi_{\theta}(a|s)}{\beta(a|s)} \hat{A}_{\theta_{\text{old}}}(s, a) \right) \quad ; \text{ Importance sampling} \\ &= \mathbb{E}_{s \sim \rho^{\pi_{\theta_{\text{old}}}}, a \sim \beta} \left[ \frac{\pi_{\theta}(a|s)}{\beta(a|s)} \hat{A}_{\theta_{\text{old}}}(s, a) \right] \end{aligned}$$

<https://lilianweng.github.io/posts/2018-04-08-policy-gradient/#trpo>

# Reinforcement Learning (GRPO)

- Sparse rewards, last step in training pipeline
- Huge dependency on the execution of the main model
- Correction terms galore, masking numerical instability
- This blog comes out...



**Your Efficient RL Framework Secretly  
Brings You Off-Policy RL Training**

Feng Yao\* Liyuan Liu\* Dinghuai Zhang Chengyu Dong Jingbo Shang Jianfeng Gao

\*: Equal Contributions (Work in Progress)

# Reinforcement Learning (numerical exactness)

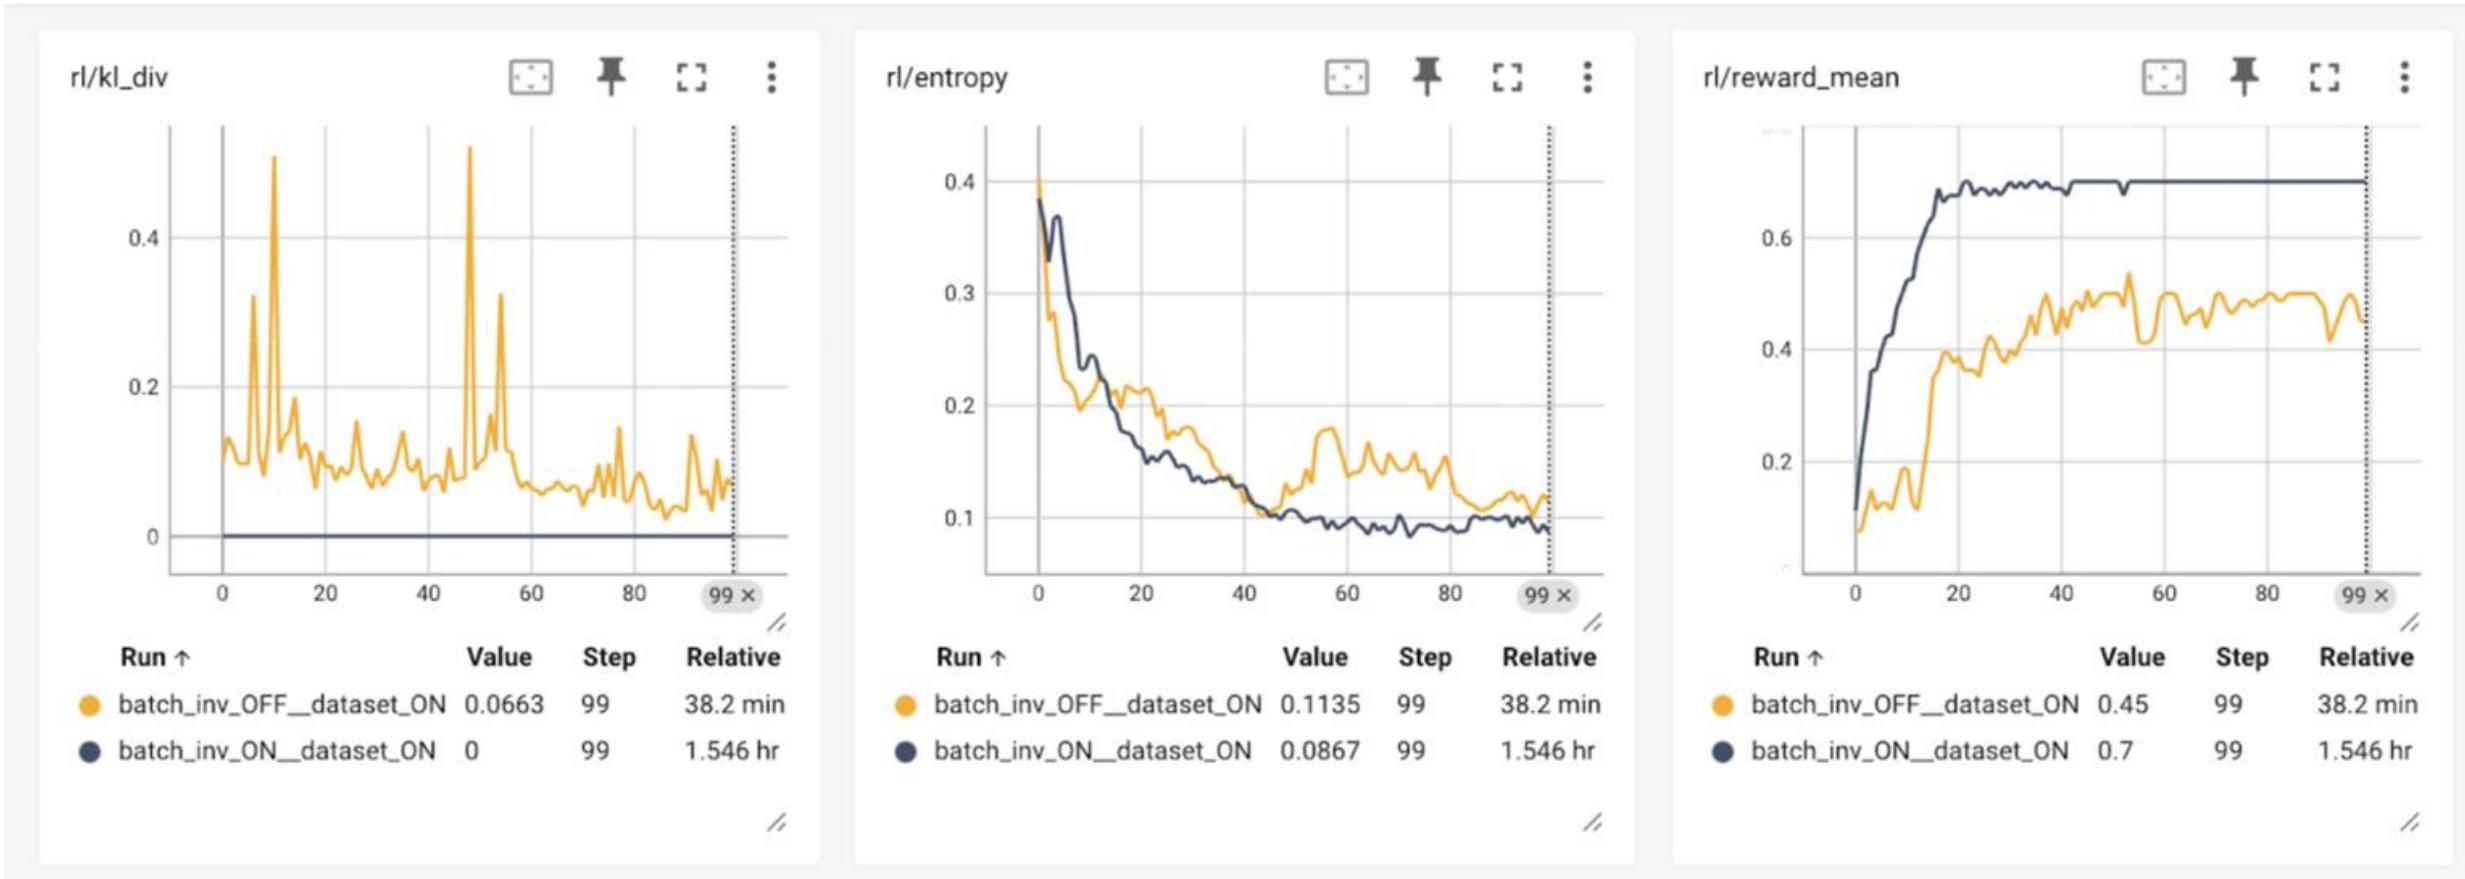
What if we ditch correction terms and use batch invariance?

```
$ VLLM_BATCH_INVARIANT=1 python simple_rl.py
Adding requests: 100%|██████████| 10/10 [00:00<00:00, 233.70it/s]
Processed prompts: 100%|██████████| 80/80 [00:02<00:00, 37.71it/s,
est. speed input: 2458.55 toks/s, output: 3770.77 toks/s]
✓ vLLM-TorchTitan bitwise determinism verified: 100 tokens match exactly

Step 479 | Loss: -0.0016 | Reward: +0.975 | Samples: 80
Sample: Natalia sold 48 clips in April. In May, she sold half as many clips as
April's ...
```

# Reinforcement Learning (numerical exactness)

Works!



## Other stability tricks / Next Steps

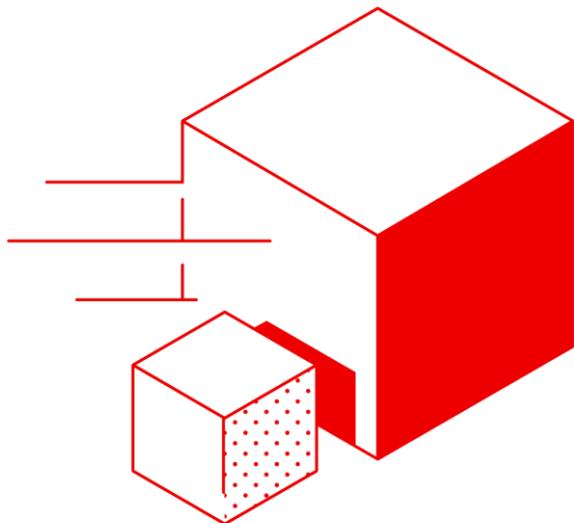
Important to note: there are other ways to stabilize RL training

- Other types of IS weights
- Use fp16 <https://arxiv.org/pdf/2510.26788>
- Rejection sampling magic

## Next Steps

- Improve performance / hardware support
- Improve compilation support (native PyTorch)
- Figure out a unification scheme for training/inference

# Get involved with the vLLM Community



## Contribute to key vLLM features

Comment and review PRs that are interesting to you. Join the discussion on RFCs. Check out "[good first issue](#)" tags.

## Give Us Feedback

We'll email you today's recording as soon as it's ready. Respond and tell us what we are doing right and what we can do better with vLLM office hours. Or comment on this slide!

## Join vLLM Developer Slack

Ask questions and engage with us via Slack. Join [here](#).

## Join Red Hat's vLLM Mission

Red Hat wants to bring open-source LLMs and vLLM to every enterprise on the planet. We are looking for vLLM Engineers to help us accomplish our mission. Apply [here](#).

# Thank you, and see you in two weeks!



**Michael Goin**

Principal Engineer, Red Hat

**vLLM Committer**



**Saša Zelenović**

Principal PMM, Red Hat



**Wentao Ye**

Machine Learning Engineer, Red Hat

**vLLM Committer**



**Bram Wasti**

Software Engineer, Meta

Version number here V00000